Personal Project_04_v10_test1_4conv-layer

April 28, 2025

```python
[1]: import pandas as pd
     import numpy as np
     import matplotlib.pyplot as plt
     %matplotlib inline
     import matplotlib.image as mpimg
     import tensorflow as tf
```

```python
[2]: # default initial values of DOE factors:
     # learning_rate = 0.001
     # dropout_value = 0.3
     # #n-conv_layers = 3
     # n_units_last_layer = 2048
     # n_filters_l1 = 32
     # n_filters_l2 = 16
```

```python
[3]: # DOE factors:
     learning_rate = 0.005
     dropout_value = 0.5
     # n-conv_layers = 4
     n_units_last_layer = 1024
     n_filters_l1 = 8
     n_filters_l2 = 64
```

```python
[4]: # other factors:
     img_size = 130
     batch_size = 32
     validation_split = 0.1  # 10% for validation
     test_split = 0.00  # 0% for testing
     shuffle_buffer_size = 1000
     seed_num = 101
     desired_accuracy = 0.99  # it should be active if EarlyStoppingCallback is␣
      ↪activated
     loss = 'binary_crossentropy'
     #optimizer = tf.keras.optimizers.RMSprop(learning_rate=learning_rate)
     optimizer = tf.keras.optimizers.Adam(learning_rate=learning_rate)
     metrics = ['accuracy']
     epochs = 15
     f_mode = 'nearest'  # fill_mode in image augmentation
```

```
My dataset_root/
   woman/
       woman_1.jpg
       woman_2.jpg
       ...
   man/
       man_1.jpg
       man_2.jpg
       ...
```

[6]:
```python
import os

DATA_DIR = "D:\\CS online courses\\Free DataSets\\Free Images\\Easier portrait␣
 ↪images_GPU_03"

# Subdirectories for each class
data_dir_woman = os.path.join(DATA_DIR, 'woman')
data_dir_man = os.path.join(DATA_DIR, 'man')

# os.listdir returns a list containing all files under the given dir
print(f"There are {len(os.listdir(data_dir_woman))} images of woman.")
print(f"There are {len(os.listdir(data_dir_man))} images of man.")
```

```
There are 471 images of woman.
There are 472 images of man.
```

[7]:
```python
image_size = (img_size, img_size)  # Resize images to this size

# Load train dataset (excluding validation & test set):
train_dataset = tf.keras.utils.image_dataset_from_directory(
    directory = DATA_DIR,
    image_size = image_size,
    batch_size = batch_size,
    label_mode='binary',
    validation_split = validation_split + test_split,  # Total split for val +␣
 ↪test
    subset = "training",
    seed = seed_num
)

# Load validation dataset
val_dataset = tf.keras.utils.image_dataset_from_directory(
    directory = DATA_DIR,
    image_size = image_size,
    batch_size = batch_size,
    label_mode='binary',
    validation_split = validation_split + test_split,
    subset = "validation",
```

```
        seed = seed_num
)

# Further manually split validation dataset to extract test dataset
val_batches = tf.data.experimental.cardinality(val_dataset)
# Compute test dataset size (number of batches)
test_size = round(val_batches.numpy() * (test_split / (validation_split +␣
  ↪test_split)))
# Split validation dataset into validation and test subsets
test_dataset = val_dataset.take(test_size)
val_dataset = val_dataset.skip(test_size)


print(f"Train batches: {tf.data.experimental.cardinality(train_dataset).
  ↪numpy()}")
print(f"Validation batches: {tf.data.experimental.cardinality(val_dataset).
  ↪numpy()}")
print(f"Test batches: {tf.data.experimental.cardinality(test_dataset).numpy()}")

# Optimize for performance
AUTOTUNE = tf.data.AUTOTUNE
training_dataset = train_dataset.cache().shuffle(shuffle_buffer_size).
  ↪prefetch(buffer_size = AUTOTUNE)
validation_dataset = val_dataset.cache().prefetch(buffer_size = AUTOTUNE)
test_dataset = test_dataset.cache().prefetch(buffer_size = AUTOTUNE)
```

```
Found 943 files belonging to 2 classes.
Using 849 files for training.
Found 943 files belonging to 2 classes.
Using 94 files for validation.
Train batches: 27
Validation batches: 3
Test batches: 0
```

[8]:
```
# Get the first batch of images and labels
for images, labels in training_dataset.take(1):
        example_batch_images = images
        example_batch_labels = labels

max_pixel = np.max(example_batch_images)
print(f"Maximum pixel value of images: {max_pixel}\n")
print(f"Shape of batch of images: {example_batch_images.shape}")
print(f"Shape of batch of labels: {example_batch_labels.shape}")
```

```
Maximum pixel value of images: 255.0

Shape of batch of images: (32, 130, 130, 3)
Shape of batch of labels: (32, 1)
```

```
[9]: '''
class EarlyStoppingCallback(tf.keras.callbacks.Callback):
    def on_epoch_end(self, epoch, logs=None):
        train_accuracy = logs.get('accuracy')
        val_accuracy = logs.get('val_accuracy')
        if train_accuracy >= desired_accuracy and val_accuracy >=⌴
  ↪desired_accuracy:
            self.model.stop_training = True
            print(f"\nReached {desired_accuracy}% accuracy so cancelling⌴
  ↪training!")
'''
```

[9]: '\nclass EarlyStoppingCallback(tf.keras.callbacks.Callback):\n   def
     on_epoch_end(self, epoch, logs=None):\n   train_accuracy =
     logs.get(\'accuracy\')\n   val_accuracy = logs.get(\'val_accuracy\')\n
     if train_accuracy >= desired_accuracy and val_accuracy >= desired_accuracy:\n
     self.model.stop_training = True\n   print(f"\nReached
     {desired_accuracy}% accuracy so cancelling training!")\n'

```
[10]: '''
from tensorflow.keras.callbacks import EarlyStopping
early_stop = EarlyStopping(monitor='val_loss', patience=3)
'''
```

[10]: "\nfrom tensorflow.keras.callbacks import EarlyStopping\nearly_stop =
      EarlyStopping(monitor='val_loss', patience=3)\n"

```
[11]: from tensorflow.keras.callbacks import LearningRateScheduler

      # Reduce LR every 10 epochs (Learning rate decay factor)
      def scheduler(epoch, lr):
          if epoch % 10 == 0 and epoch > 0:
              return lr * 1.0
          return lr

      lr_callback = LearningRateScheduler(scheduler)
```

```
[12]: # augmentation_model
      def augment_model():
          """Creates a model (layers stacked on top of each other) for augmenting⌴
        ↪images of woman and man.

          Returns:
              tf.keras.Model: The model made up of the layers that will be used to⌴
        ↪augment the images of woman and man.
          """
```

```
    augmentation_model = tf.keras.Sequential([
        # Specify the input shape.
        tf.keras.Input(shape = (img_size, img_size, 3)),

        tf.keras.layers.RandomFlip("horizontal"),
        tf.keras.layers.RandomRotation(0.1, fill_mode = f_mode),
        #tf.keras.layers.RandomTranslation(0.1, 0.1, fill_mode = f_mode),
        #tf.keras.layers.RandomZoom(0.1, fill_mode=f_mode)
        ])

    return augmentation_model
```

```
[13]: def create_and_compile_model():
    """Creates, compiles and trains the model to predict woman and man images.

    Returns:
        tf.keras.Model: The model that will be trained to predict woman and man⊔
    ↪images.
    """

    augmentation_layers = augment_model()

    model = tf.keras.Sequential([
        # Note: the input shape is the desired size of the image: 150x150 with⊔
    ↪3 bytes for color
        tf.keras.layers.InputLayer(shape = (img_size, img_size, 3)),
        augmentation_layers,
        tf.keras.layers.Rescaling(1./255),
        #####    CONV_LAYER_1:    #####
        tf.keras.layers.Conv2D(n_filters_l1, (4, 4), activation = 'linear'),
        tf.keras.layers.MaxPooling2D(2, 2),
        #####    CONV_LAYER_2:    #####
        tf.keras.layers.Conv2D(n_filters_l2, (3, 3), activation = 'relu'),
        tf.keras.layers.MaxPooling2D(2, 2),
        #####    CONV_LAYER_3:    #####
        tf.keras.layers.Conv2D(64, (3, 3), activation = 'relu'),
        tf.keras.layers.MaxPooling2D(2, 2),
        #####    CONV_LAYER_4:    #####
        tf.keras.layers.Conv2D(64, (3, 3), activation = 'relu'),
        tf.keras.layers.MaxPooling2D(2, 2),
        tf.keras.layers.Flatten(),
        tf.keras.layers.Dropout(dropout_value),
        #####    BEFORE_LAST_LAYER:    #####
        tf.keras.layers.Dense(n_units_last_layer, activation = 'relu'),
        # It will contain a value from 0-1 where 0 for the class 'female' and 1⊔
    ↪for the 'male'
        tf.keras.layers.Dense(1, activation = 'sigmoid')])
```

```
    model.compile(
        loss = loss,
        optimizer = optimizer,
        metrics = metrics
    )

    return model
```

[14]:
```
# Create the compiled but untrained model
model = create_and_compile_model()
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
sequential (Sequential)	(None, 130, 130, 3)	0
rescaling (Rescaling)	(None, 130, 130, 3)	0
conv2d (Conv2D)	(None, 127, 127, 8)	392
max_pooling2d (MaxPooling2D)	(None, 63, 63, 8)	0
conv2d_1 (Conv2D)	(None, 61, 61, 64)	4,672
max_pooling2d_1 (MaxPooling2D)	(None, 30, 30, 64)	0
conv2d_2 (Conv2D)	(None, 28, 28, 64)	36,928
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 64)	0
conv2d_3 (Conv2D)	(None, 12, 12, 64)	36,928
max_pooling2d_3 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten (Flatten)	(None, 2304)	0
dropout (Dropout)	(None, 2304)	0
dense (Dense)	(None, 1024)	2,360,320
dense_1 (Dense)	(None, 1)	1,025

```
Total params: 2,440,265 (9.31 MB)

Trainable params: 2,440,265 (9.31 MB)

Non-trainable params: 0 (0.00 B)
```

[15]:
```python
'''
training_history = model.fit(
    training_dataset,
    epochs = epochs,
    validation_data = validation_dataset,
    callbacks = [EarlyStoppingCallback()],
    verbose = 2
)
'''
```

[15]: '\ntraining_history = model.fit(\n   training_dataset,\n   epochs = epochs,\n validation_data = validation_dataset,\n   callbacks = [EarlyStoppingCallback()],\n   verbose = 2\n)\n'

[16]:
```python
'''
training_history = model.fit(
    training_dataset,
    epochs = epochs,
    validation_data = validation_dataset,
    callbacks=[early_stop],
    verbose = 2
)
'''
```

[16]: '\ntraining_history = model.fit(\n   training_dataset,\n   epochs = epochs,\n validation_data = validation_dataset,\n   callbacks=[early_stop],\n   verbose = 2\n)\n'

[17]:
```python
training_history = model.fit(
    training_dataset,
    epochs = epochs,
    validation_data = validation_dataset,
    callbacks = [lr_callback],
    verbose = 2
)
```

```
Epoch 1/15
27/27 - 4s - 137ms/step - accuracy: 0.5300 - loss: 0.9119 - val_accuracy: 0.6596
- val_loss: 0.6694 - learning_rate: 0.0050
Epoch 2/15
```

```
27/27 - 2s - 60ms/step - accuracy: 0.5453 - loss: 0.6855 - val_accuracy: 0.6809
- val_loss: 0.6395 - learning_rate: 0.0050
Epoch 3/15
27/27 - 2s - 56ms/step - accuracy: 0.6337 - loss: 0.6470 - val_accuracy: 0.7128
- val_loss: 0.5766 - learning_rate: 0.0050
Epoch 4/15
27/27 - 2s - 61ms/step - accuracy: 0.6561 - loss: 0.6207 - val_accuracy: 0.7234
- val_loss: 0.5846 - learning_rate: 0.0050
Epoch 5/15
27/27 - 2s - 61ms/step - accuracy: 0.5807 - loss: 0.6576 - val_accuracy: 0.7447
- val_loss: 0.5829 - learning_rate: 0.0050
Epoch 6/15
27/27 - 2s - 58ms/step - accuracy: 0.6596 - loss: 0.6392 - val_accuracy: 0.6809
- val_loss: 0.6260 - learning_rate: 0.0050
Epoch 7/15
27/27 - 2s - 58ms/step - accuracy: 0.6784 - loss: 0.6063 - val_accuracy: 0.6809
- val_loss: 0.6325 - learning_rate: 0.0050
Epoch 8/15
27/27 - 2s - 58ms/step - accuracy: 0.6678 - loss: 0.6050 - val_accuracy: 0.6915
- val_loss: 0.5569 - learning_rate: 0.0050
Epoch 9/15
27/27 - 2s - 59ms/step - accuracy: 0.6655 - loss: 0.6061 - val_accuracy: 0.6596
- val_loss: 0.6335 - learning_rate: 0.0050
Epoch 10/15
27/27 - 2s - 57ms/step - accuracy: 0.6749 - loss: 0.6015 - val_accuracy: 0.6809
- val_loss: 0.5832 - learning_rate: 0.0050
Epoch 11/15
27/27 - 2s - 62ms/step - accuracy: 0.7173 - loss: 0.5628 - val_accuracy: 0.7447
- val_loss: 0.5172 - learning_rate: 0.0050
Epoch 12/15
27/27 - 2s - 60ms/step - accuracy: 0.7020 - loss: 0.5821 - val_accuracy: 0.7234
- val_loss: 0.5832 - learning_rate: 0.0050
Epoch 13/15
27/27 - 2s - 59ms/step - accuracy: 0.7079 - loss: 0.5701 - val_accuracy: 0.7340
- val_loss: 0.5606 - learning_rate: 0.0050
Epoch 14/15
27/27 - 2s - 62ms/step - accuracy: 0.6902 - loss: 0.5738 - val_accuracy: 0.7234
- val_loss: 0.5693 - learning_rate: 0.0050
Epoch 15/15
27/27 - 2s - 61ms/step - accuracy: 0.7020 - loss: 0.5661 - val_accuracy: 0.7340
- val_loss: 0.5579 - learning_rate: 0.0050
```

[18]:
```python
#from tensorflow.keras.models import load_model
#model.save('gender_recognition_project04_v10.h5')
```

[19]:
```python
model.metrics_names
```

```
[19]: ['loss', 'compile_metrics']
```

```
[20]: result_history = pd.DataFrame(model.history.history)
      result_history.head(15)
```

```
[20]:      accuracy      loss  val_accuracy  val_loss  learning_rate
      0    0.530035  0.911853      0.659574  0.669449          0.005
      1    0.545347  0.685545      0.680851  0.639519          0.005
      2    0.633687  0.646964      0.712766  0.576611          0.005
      3    0.656066  0.620689      0.723404  0.584580          0.005
      4    0.580683  0.657592      0.744681  0.582932          0.005
      5    0.659600  0.639182      0.680851  0.626032          0.005
      6    0.678445  0.606348      0.680851  0.632506          0.005
      7    0.667845  0.604961      0.691489  0.556852          0.005
      8    0.665489  0.606149      0.659574  0.633534          0.005
      9    0.674912  0.601478      0.680851  0.583224          0.005
      10   0.717314  0.562813      0.744681  0.517247          0.005
      11   0.702002  0.582057      0.723404  0.583173          0.005
      12   0.707892  0.570093      0.734043  0.560596          0.005
      13   0.690224  0.573779      0.723404  0.569268          0.005
      14   0.702002  0.566097      0.734043  0.557947          0.005
```

```
[21]: result_history[['loss', 'val_loss']].plot(figsize=(5, 3))
```

```
[21]: <Axes: >
```



```
[22]: result_history[['accuracy', 'val_accuracy']].plot(figsize=(5, 3))
```

```
[22]: <Axes: >
```

```
[23]: print(model.metrics_names)
      print(model.evaluate(validation_dataset))
```

```
['loss', 'compile_metrics']
3/3                0s 18ms/step -
accuracy: 0.7264 - loss: 0.5698
[0.5579467415809631, 0.7340425252914429]
```

```
[24]: from sklearn.metrics import classification_report, confusion_matrix

      y_true = np.concatenate([y.numpy() for _, y in validation_dataset])
      y_pred_prob = model.predict(validation_dataset)
      # Convert probabilities to class labels (0:Female or 1:Male)
      y_pred = (y_pred_prob > 0.5).astype(int).flatten()

      print("Classification Report:\n", classification_report(y_true, y_pred,␣
       ↪target_names=['Female', 'Male']))
```

```
3/3                0s 57ms/step
Classification Report:
               precision    recall  f1-score   support

      Female       0.67      0.78      0.72        41
        Male       0.80      0.70      0.75        53

    accuracy                           0.73        94
   macro avg       0.74      0.74      0.73        94
weighted avg       0.74      0.73      0.74        94
```

10

```python
[25]: import tensorflow as tf
      import numpy as np
      import matplotlib.pyplot as plt
      from tensorflow.keras.models import Model
      from tensorflow.keras.utils import load_img, img_to_array

      img_size = img_size
      model = tf.keras.models.load_model("gender_recognition_project04_v10.h5")

      # Load your personal image if you are interested to predict:
      your_image_path = "D:\\Hossein's desktop files in Microsoft Studio␣
        ↪Laptop\\Personal Photos\\Hossein_10.jpg"

      img = load_img(your_image_path, target_size=(img_size, img_size))
      final_img = img_to_array(img)
      # Adding a batch dimension:
      final_img = np.expand_dims(final_img, axis=0)
      prediction = model.predict(final_img)
      result = "Female" if prediction > 0.5 else "Male"
      if result=="Female":
          confidence = (model.predict(final_img)[0][0])*100
      else:
          confidence = (1-model.predict(final_img)[0][0])*100
      print(f"Prediction result: {result} (confidence= {confidence:.2f} %)")

      # Visualize CNN Layers
      successive_feature_maps = visualization_model.predict(final_img)
      layer_names = [layer.name for layer in model.layers]

      for layer_name, feature_map in zip(layer_names, successive_feature_maps):
          if len(feature_map.shape) == 4:  # Only visualize conv/maxpool layers
              n_features = feature_map.shape[-1]  # Number of filters
              size = feature_map.shape[1]  # Feature map size
              display_grid = np.zeros((size, size * n_features))

              for i in range(n_features):
                  x = feature_map[0, :, :, i]
                  x -= x.mean()
                  x /= (x.std() + 1e-8)  # Normalize
                  x *= 64
                  x += 128
                  x = np.clip(x, 0, 255).astype('uint8')  # Convert to image format
                  display_grid[:, i * size: (i + 1) * size] = x

              scale = 20. / n_features
              plt.figure(figsize=(scale * n_features, scale))
              plt.title(layer_name)
```

11

```
        plt.grid(False)
        plt.imshow(display_grid, aspect='auto', cmap='cividis')
        plt.show()
```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.

```
1/1                  0s 109ms/step
1/1                  0s 39ms/step
Prediction result: Male (confidence= 94.19 %)
```

```
---------------------------------------------------------------------------
NameError                                 Traceback (most recent call last)
Cell In[25], line 26
     23 print(f"Prediction result: {result} (confidence= {confidence:.2f} %)")
     25 # Visualize CNN Layers
---> 26 successive_feature_maps = visualization_model.predict(final_img)
     27 layer_names = [layer.name for layer in model.layers]
     29 for layer_name, feature_map in zip(layer_names, successive_feature_maps :

NameError: name 'visualization_model' is not defined
```

[ ]:

[ ]:

[ ]: