

Personal Project_04_v10_test1_2conv-layer_run042_advanced contro_autorun

May 5, 2025

```
[1]: from tensorflow.keras.callbacks import LearningRateScheduler
from sklearn.metrics import classification_report, confusion_matrix
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import matplotlib.image as mpimg
import tensorflow as tf
import os

ACC=0.1
try_num = 1

while (ACC<0.88 and try_num<10):
    # DOE factors:
    learning_rate = 0.005
    dropout_value = 0.5
    # n_conv_layers = 2
    n_units_last_layer = 1024
    n_filters_l1 = 8
    n_filters_l2 = 16

    # other factors:
    img_size = 130
    batch_size = 32
    validation_split = 0.1 # 10% for validation
    test_split = 0.00 # 0% for testing
    shuffle_buffer_size = 1000
    seed_num = 101
    desired_accuracy = 0.99 # it should be active if EarlyStoppingCallback is
    ↪activated

    loss = 'binary_crossentropy'
    #optimizer = tf.keras.optimizers.RMSprop(learning_rate=learning_rate)
    optimizer = tf.keras.optimizers.Adam(learning_rate=learning_rate)
    metrics = ['accuracy']
```

```

epochs = 29
f_mode = 'nearest' # fill_mode in image augmentation

DATA_DIR = "D:\\CS online courses\\Free DataSets\\Free Images\\Easier_
↳portrait images_GPU_03"
#DATA_DIR = "/Users/hosseini/Downloads/Easier portrait images_GPU_03"

# Subdirectories for each class
data_dir_woman = os.path.join(DATA_DIR, 'woman')
data_dir_man = os.path.join(DATA_DIR, 'man')
image_size = (img_size, img_size) # Resize images to this size
# Load train dataset (excluding validation & test set):
train_dataset = tf.keras.utils.image_dataset_from_directory(
    directory = DATA_DIR,
    image_size = image_size,
    batch_size = batch_size,
    label_mode='binary',
    validation_split = validation_split + test_split, # Total split for_
↳val + test
    subset = "training",
    seed = seed_num
)
# Load validation dataset
val_dataset = tf.keras.utils.image_dataset_from_directory(
    directory = DATA_DIR,
    image_size = image_size,
    batch_size = batch_size,
    label_mode='binary',
    validation_split = validation_split + test_split,
    subset = "validation",
    seed = seed_num
)
# Further manually split validation dataset to extract test dataset
val_batches = tf.data.experimental.cardinality(val_dataset)
# Compute test dataset size (number of batches)
test_size = round(val_batches.numpy() * (test_split / (validation_split +_
↳test_split)))
# Split validation dataset into validation and test subsets
test_dataset = val_dataset.take(test_size)
val_dataset = val_dataset.skip(test_size)
# Optimize for performance
AUTOTUNE = tf.data.AUTOTUNE
training_dataset = train_dataset.cache().shuffle(shuffle_buffer_size).
↳prefetch(buffer_size = AUTOTUNE)
validation_dataset = val_dataset.cache().prefetch(buffer_size = AUTOTUNE)
test_dataset = test_dataset.cache().prefetch(buffer_size = AUTOTUNE)

```

```

# Get the first batch of images and labels
for images, labels in training_dataset.take(1):
    example_batch_images = images
    example_batch_labels = labels
max_pixel = np.max(example_batch_images)

# Reduce LR every 10 epochs (Learning rate decay factor)
def scheduler(epoch, lr):
    if epoch < 10:
        if epoch % 5 == 0 and epoch > 0:
            return lr / 1
        return lr
    elif epoch < 15:
        if epoch % 5 == 0 and epoch > 0:
            return lr / 1
        return lr
    elif epoch < 30:
        if epoch % 5 == 0 and epoch > 0:
            return lr / 1
        return lr
    else:
        return lr
lr_callback = LearningRateScheduler(scheduler)

# augmentation_model
def augment_model():
    augmentation_model = tf.keras.Sequential([
        # Specify the input shape.
        tf.keras.Input(shape = (img_size, img_size, 3)),

        tf.keras.layers.RandomFlip("horizontal"),
        tf.keras.layers.RandomRotation(0.1, fill_mode = f_mode),
        #tf.keras.layers.RandomTranslation(0.1, 0.1, fill_mode = f_mode),
        #tf.keras.layers.RandomZoom(0.1, fill_mode=f_mode)
    ])
    return augmentation_model

def create_and_compile_model():
    augmentation_layers = augment_model()
    model = tf.keras.Sequential([
        # Note: the input shape is the desired size of the image: 150x150
        ↪with 3 bytes for color
        tf.keras.layers.InputLayer(shape = (img_size, img_size, 3)),
        augmentation_layers,
        tf.keras.layers.Rescaling(1./255),
        ##### CONV_LAYER_1: #####
        tf.keras.layers.Conv2D(n_filters_l1, (4, 4), activation = 'linear'),

```

```

        tf.keras.layers.MaxPooling2D(2, 2),
        ##### CONV_LAYER_2: #####
        tf.keras.layers.Conv2D(n_filters_12, (3, 3), activation = 'relu'),
        tf.keras.layers.MaxPooling2D(2, 2),
        tf.keras.layers.Flatten(),
        tf.keras.layers.Dropout(dropout_value),
        ##### BEFORE_LAST_LAYER: #####
        tf.keras.layers.Dense(n_units_last_layer, activation = 'relu'),
        # It will contain a value from 0-1 where 0 for the class 'female'
        ↪ and 1 for the 'male'
        tf.keras.layers.Dense(1, activation = 'sigmoid'))])
    model.compile(
        loss = loss,
        optimizer = optimizer,
        metrics = metrics
    )
    return model

# Create the compiled but untrained model
def reset_weights(model):
    for layer in model.layers:
        if hasattr(layer, 'kernel_initializer'):
            layer.kernel.assign(layer.kernel_initializer(layer.kernel.
            ↪ shape))
        if hasattr(layer, 'bias_initializer'):
            layer.bias.assign(layer.bias_initializer(layer.bias.shape))

model = create_and_compile_model()
reset_weights(model) # Reset all layer weights
training_history = model.fit(training_dataset,
                             epochs=epochs,
                             validation_data=validation_dataset,
                             callbacks=[lr_callback],
                             verbose=2)

result_history = pd.DataFrame(model.history.history)
ACC = result_history['val_accuracy'].iloc[-1]
print(f"Current validation accuracy: {ACC}")
model.save('trained_model_run42_advanced_control.h5')
# Restart script
print("Resetting all weights...")
print(f'Current number of trials: {try_num}')
try_num += 1
result_history[['loss', 'val_loss']].plot(figsize=(5, 3))
result_history[['accuracy', 'val_accuracy']].plot(figsize=(5, 3))
plt.show()
print(model.metrics_names)
print(model.evaluate(validation_dataset))

```

```

y_true = np.concatenate([y.numpy() for _, y in validation_dataset])
y_pred_prob = model.predict(validation_dataset)
# Convert probabilities to class labels (0:Female or 1:Male)
y_pred = (y_pred_prob > 0.5).astype(int).flatten()
print("Classification Report:\n", classification_report(y_true, y_pred,
↳target_names=['Female', 'Male']))

result_history.head(15)

```

Found 943 files belonging to 2 classes.

Using 849 files for training.

Found 943 files belonging to 2 classes.

Using 94 files for validation.

Epoch 1/29

27/27 - 4s - 137ms/step - accuracy: 0.5241 - loss: 2.9836 - val_accuracy: 0.5851
- val_loss: 0.6829 - learning_rate: 0.0050

Epoch 2/29

27/27 - 2s - 70ms/step - accuracy: 0.6019 - loss: 0.6791 - val_accuracy: 0.5851
- val_loss: 0.6669 - learning_rate: 0.0050

Epoch 3/29

27/27 - 2s - 68ms/step - accuracy: 0.5901 - loss: 0.6753 - val_accuracy: 0.6702
- val_loss: 0.6742 - learning_rate: 0.0050

Epoch 4/29

27/27 - 2s - 69ms/step - accuracy: 0.6419 - loss: 0.6400 - val_accuracy: 0.6277
- val_loss: 0.6250 - learning_rate: 0.0050

Epoch 5/29

27/27 - 2s - 78ms/step - accuracy: 0.6678 - loss: 0.6058 - val_accuracy: 0.7340
- val_loss: 0.5599 - learning_rate: 0.0050

Epoch 6/29

27/27 - 2s - 72ms/step - accuracy: 0.6855 - loss: 0.6090 - val_accuracy: 0.7340
- val_loss: 0.6086 - learning_rate: 0.0050

Epoch 7/29

27/27 - 2s - 73ms/step - accuracy: 0.7102 - loss: 0.5634 - val_accuracy: 0.7766
- val_loss: 0.5745 - learning_rate: 0.0050

Epoch 8/29

27/27 - 2s - 76ms/step - accuracy: 0.7244 - loss: 0.5360 - val_accuracy: 0.7660
- val_loss: 0.5326 - learning_rate: 0.0050

Epoch 9/29

27/27 - 2s - 67ms/step - accuracy: 0.7256 - loss: 0.5567 - val_accuracy: 0.7766
- val_loss: 0.5552 - learning_rate: 0.0050

Epoch 10/29

27/27 - 2s - 70ms/step - accuracy: 0.7397 - loss: 0.5029 - val_accuracy: 0.7340
- val_loss: 0.5408 - learning_rate: 0.0050

Epoch 11/29

27/27 - 2s - 73ms/step - accuracy: 0.7703 - loss: 0.4850 - val_accuracy: 0.7553
- val_loss: 0.5296 - learning_rate: 0.0050

Epoch 12/29

27/27 - 2s - 84ms/step - accuracy: 0.7951 - loss: 0.4414 - val_accuracy: 0.6702

- val_loss: 0.7635 - learning_rate: 0.0050
Epoch 13/29
27/27 - 2s - 75ms/step - accuracy: 0.8080 - loss: 0.4288 - val_accuracy: 0.8085
- val_loss: 0.4418 - learning_rate: 0.0050
Epoch 14/29
27/27 - 2s - 72ms/step - accuracy: 0.7833 - loss: 0.4486 - val_accuracy: 0.7553
- val_loss: 0.5014 - learning_rate: 0.0050
Epoch 15/29
27/27 - 2s - 72ms/step - accuracy: 0.8115 - loss: 0.4241 - val_accuracy: 0.7766
- val_loss: 0.4649 - learning_rate: 0.0050
Epoch 16/29
27/27 - 2s - 68ms/step - accuracy: 0.8068 - loss: 0.4465 - val_accuracy: 0.7979
- val_loss: 0.4181 - learning_rate: 0.0050
Epoch 17/29
27/27 - 2s - 67ms/step - accuracy: 0.8221 - loss: 0.3806 - val_accuracy: 0.7766
- val_loss: 0.5044 - learning_rate: 0.0050
Epoch 18/29
27/27 - 2s - 68ms/step - accuracy: 0.7915 - loss: 0.4238 - val_accuracy: 0.7979
- val_loss: 0.4731 - learning_rate: 0.0050
Epoch 19/29
27/27 - 2s - 67ms/step - accuracy: 0.8221 - loss: 0.3890 - val_accuracy: 0.8085
- val_loss: 0.4946 - learning_rate: 0.0050
Epoch 20/29
27/27 - 2s - 67ms/step - accuracy: 0.8375 - loss: 0.3886 - val_accuracy: 0.8404
- val_loss: 0.3906 - learning_rate: 0.0050
Epoch 21/29
27/27 - 2s - 67ms/step - accuracy: 0.8492 - loss: 0.3548 - val_accuracy: 0.7979
- val_loss: 0.4425 - learning_rate: 0.0050
Epoch 22/29
27/27 - 2s - 67ms/step - accuracy: 0.8445 - loss: 0.3464 - val_accuracy: 0.8298
- val_loss: 0.4605 - learning_rate: 0.0050
Epoch 23/29
27/27 - 2s - 67ms/step - accuracy: 0.8610 - loss: 0.3066 - val_accuracy: 0.8298
- val_loss: 0.4365 - learning_rate: 0.0050
Epoch 24/29
27/27 - 2s - 68ms/step - accuracy: 0.8763 - loss: 0.2850 - val_accuracy: 0.8085
- val_loss: 0.5054 - learning_rate: 0.0050
Epoch 25/29
27/27 - 2s - 67ms/step - accuracy: 0.8645 - loss: 0.3144 - val_accuracy: 0.8085
- val_loss: 0.4857 - learning_rate: 0.0050
Epoch 26/29
27/27 - 2s - 67ms/step - accuracy: 0.8634 - loss: 0.3359 - val_accuracy: 0.8404
- val_loss: 0.3977 - learning_rate: 0.0050
Epoch 27/29
27/27 - 2s - 68ms/step - accuracy: 0.8846 - loss: 0.2822 - val_accuracy: 0.8298
- val_loss: 0.4897 - learning_rate: 0.0050
Epoch 28/29
27/27 - 2s - 67ms/step - accuracy: 0.8799 - loss: 0.2776 - val_accuracy: 0.7872

- val_loss: 0.5563 - learning_rate: 0.0050

Epoch 29/29

27/27 - 2s - 67ms/step - accuracy: 0.8622 - loss: 0.3202 - val_accuracy: 0.8298



- val_loss: 0.4403 - learning_rate: 0.0050

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

Current validation accuracy: 0.8297872543334961

Resetting all weights...

Current number of trials: 1



```
['loss', 'compile_metrics']
3/3          0s 17ms/step -
accuracy: 0.8172 - loss: 0.4643
[0.4402659833431244, 0.8297872543334961]
3/3          0s 49ms/step
Classification Report:
              precision    recall  f1-score   support

   Female           0.79       0.83       0.81         41
    Male           0.86       0.83       0.85         53

 accuracy                   0.83         94
  macro avg           0.83       0.83       0.83         94
 weighted avg          0.83       0.83       0.83         94
```

Found 943 files belonging to 2 classes.

Using 849 files for training.

Found 943 files belonging to 2 classes.

Using 94 files for validation.

Epoch 1/29

27/27 - 4s - 139ms/step - accuracy: 0.6478 - loss: 1.9918 - val_accuracy: 0.7234
- val_loss: 0.5742 - learning_rate: 0.0050

Epoch 2/29

27/27 - 2s - 68ms/step - accuracy: 0.7102 - loss: 0.5456 - val_accuracy: 0.7872
- val_loss: 0.5530 - learning_rate: 0.0050

Epoch 3/29

27/27 - 2s - 67ms/step - accuracy: 0.7044 - loss: 0.5561 - val_accuracy: 0.7447
- val_loss: 0.5362 - learning_rate: 0.0050

Epoch 4/29

27/27 - 2s - 67ms/step - accuracy: 0.7562 - loss: 0.5157 - val_accuracy: 0.7553
- val_loss: 0.4819 - learning_rate: 0.0050

Epoch 5/29

27/27 - 2s - 67ms/step - accuracy: 0.7550 - loss: 0.5572 - val_accuracy: 0.7234
- val_loss: 0.5941 - learning_rate: 0.0050

Epoch 6/29

27/27 - 2s - 67ms/step - accuracy: 0.7385 - loss: 0.5349 - val_accuracy: 0.7979
- val_loss: 0.6125 - learning_rate: 0.0050

Epoch 7/29

27/27 - 2s - 66ms/step - accuracy: 0.7621 - loss: 0.4866 - val_accuracy: 0.7766
- val_loss: 0.5653 - learning_rate: 0.0050

Epoch 8/29

27/27 - 2s - 67ms/step - accuracy: 0.7715 - loss: 0.4939 - val_accuracy: 0.8085
- val_loss: 0.4927 - learning_rate: 0.0050

Epoch 9/29

27/27 - 2s - 67ms/step - accuracy: 0.7880 - loss: 0.4740 - val_accuracy: 0.7979
- val_loss: 0.4676 - learning_rate: 0.0050

Epoch 10/29
 27/27 - 2s - 68ms/step - accuracy: 0.8045 - loss: 0.4564 - val_accuracy: 0.7872
 - val_loss: 0.5353 - learning_rate: 0.0050

Epoch 11/29
 27/27 - 2s - 67ms/step - accuracy: 0.8021 - loss: 0.4411 - val_accuracy: 0.8191
 - val_loss: 0.4392 - learning_rate: 0.0050

Epoch 12/29
 27/27 - 2s - 67ms/step - accuracy: 0.8233 - loss: 0.4097 - val_accuracy: 0.7979
 - val_loss: 0.4771 - learning_rate: 0.0050

Epoch 13/29
 27/27 - 2s - 67ms/step - accuracy: 0.8292 - loss: 0.3962 - val_accuracy: 0.8404
 - val_loss: 0.3880 - learning_rate: 0.0050

Epoch 14/29
 27/27 - 2s - 67ms/step - accuracy: 0.8080 - loss: 0.4072 - val_accuracy: 0.8298
 - val_loss: 0.4432 - learning_rate: 0.0050

Epoch 15/29
 27/27 - 2s - 67ms/step - accuracy: 0.8280 - loss: 0.3881 - val_accuracy: 0.8617
 - val_loss: 0.3907 - learning_rate: 0.0050

Epoch 16/29
 27/27 - 2s - 67ms/step - accuracy: 0.8327 - loss: 0.3737 - val_accuracy: 0.8298
 - val_loss: 0.5187 - learning_rate: 0.0050

Epoch 17/29
 27/27 - 2s - 67ms/step - accuracy: 0.8481 - loss: 0.3500 - val_accuracy: 0.8191
 - val_loss: 0.4510 - learning_rate: 0.0050

Epoch 18/29
 27/27 - 2s - 67ms/step - accuracy: 0.8386 - loss: 0.3600 - val_accuracy: 0.8298
 - val_loss: 0.5003 - learning_rate: 0.0050

Epoch 19/29
 27/27 - 2s - 67ms/step - accuracy: 0.8481 - loss: 0.3481 - val_accuracy: 0.8404
 - val_loss: 0.3387 - learning_rate: 0.0050

Epoch 20/29
 27/27 - 2s - 67ms/step - accuracy: 0.8563 - loss: 0.3295 - val_accuracy: 0.8404
 - val_loss: 0.4390 - learning_rate: 0.0050

Epoch 21/29
 27/27 - 2s - 67ms/step - accuracy: 0.8457 - loss: 0.3594 - val_accuracy: 0.8404
 - val_loss: 0.4201 - learning_rate: 0.0050

Epoch 22/29
 27/27 - 2s - 67ms/step - accuracy: 0.8363 - loss: 0.3688 - val_accuracy: 0.7872
 - val_loss: 0.5267 - learning_rate: 0.0050

Epoch 23/29
 27/27 - 2s - 67ms/step - accuracy: 0.8386 - loss: 0.3630 - val_accuracy: 0.8085
 - val_loss: 0.4585 - learning_rate: 0.0050

Epoch 24/29
 27/27 - 2s - 67ms/step - accuracy: 0.8787 - loss: 0.2830 - val_accuracy: 0.8191
 - val_loss: 0.5042 - learning_rate: 0.0050

Epoch 25/29
 27/27 - 2s - 67ms/step - accuracy: 0.8716 - loss: 0.2890 - val_accuracy: 0.8617
 - val_loss: 0.5012 - learning_rate: 0.0050

Epoch 26/29

27/27 - 2s - 67ms/step - accuracy: 0.8704 - loss: 0.3109 - val_accuracy: 0.8298
- val_loss: 0.4689 - learning_rate: 0.0050

Epoch 27/29

27/27 - 2s - 67ms/step - accuracy: 0.8787 - loss: 0.2732 - val_accuracy: 0.8085
- val_loss: 0.4356 - learning_rate: 0.0050

Epoch 28/29

27/27 - 2s - 68ms/step - accuracy: 0.9011 - loss: 0.2619 - val_accuracy: 0.8191
- val_loss: 0.5356 - learning_rate: 0.0050

Epoch 29/29

27/27 - 2s - 67ms/step - accuracy: 0.8846 - loss: 0.2646 - val_accuracy: 0.8830
- val_loss: 0.3757 - learning_rate: 0.0050


WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.


`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

Current validation accuracy: 0.8829787373542786

Resetting all weights...

Current number of trials: 2





```
['loss', 'compile_metrics']
3/3          0s 17ms/step -
accuracy: 0.8829 - loss: 0.3935
[0.37573546171188354, 0.8829787373542786]
3/3          0s 45ms/step
Classification Report:
              precision    recall  f1-score   support

   Female       0.83        0.93        0.87         41
    Male       0.94        0.85        0.89         53

 accuracy                   0.88         94
 macro avg       0.88        0.89        0.88         94
weighted avg       0.89        0.88        0.88         94
```

```
[1]: accuracy      loss  val_accuracy  val_loss  learning_rate
0    0.647821  1.991831    0.723404  0.574178      0.005
1    0.710247  0.545627    0.787234  0.552998      0.005
2    0.704358  0.556103    0.744681  0.536161      0.005
3    0.756184  0.515687    0.755319  0.481906      0.005
4    0.755006  0.557162    0.723404  0.594113      0.005
5    0.738516  0.534944    0.797872  0.612482      0.005
6    0.762073  0.486636    0.776596  0.565275      0.005
7    0.771496  0.493853    0.808511  0.492743      0.005
8    0.787986  0.473971    0.797872  0.467577      0.005
9    0.804476  0.456402    0.787234  0.535287      0.005
10   0.802120  0.441130    0.819149  0.439167      0.005
11   0.823322  0.409732    0.797872  0.477067      0.005
12   0.829211  0.396230    0.840426  0.388031      0.005
```

13	0.808009	0.407230	0.829787	0.443219	0.005
14	0.828033	0.388145	0.861702	0.390725	0.005

[]:

[]:

[]: