

# Personal Project\_04\_v10\_test1\_5conv-layer-Copy1

April 30, 2025

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import matplotlib.image as mpimg
import tensorflow as tf
```

```
[2]: # default initial values of DOE factors:
# learning_rate = 0.001
# dropout_value = 0.3
# #n-conv_layers = 3
# n_units_last_layer = 2048
# n_filters_l1 = 32
# n_filters_l2 = 16
```

```
[3]: # DOE factors:
learning_rate = 0.001
dropout_value = 0.3
# n-conv_layers = 5
n_units_last_layer = 2048
n_filters_l1 = 16
n_filters_l2 = 32
```

```
[4]: # other factors:
img_size = 130
batch_size = 32
validation_split = 0.1 # 10% for validation
test_split = 0.00 # 0% for testing
shuffle_buffer_size = 1000
seed_num = 101
desired_accuracy = 0.99 # it should be active if EarlyStoppingCallback is
↳activated
loss = 'binary_crossentropy'
#optimizer = tf.keras.optimizers.RMSprop(learning_rate=learning_rate)
optimizer = tf.keras.optimizers.Adam(learning_rate=learning_rate)
metrics = ['accuracy']
epochs = 100
f_mode = 'nearest' # fill_mode in image augmentation
```

```

My dataset_root/
  woman/
    woman_1.jpg
    woman_2.jpg
    ...
  man/
    man_1.jpg
    man_2.jpg
    ...

```

```

[6]: import os

DATA_DIR = "D:\\CS online courses\\Free DataSets\\Free Images\\Easier portrait_
↳images_GPU_03"

# Subdirectories for each class
data_dir_woman = os.path.join(DATA_DIR, 'woman')
data_dir_man = os.path.join(DATA_DIR, 'man')

# os.listdir returns a list containing all files under the given dir
print(f"There are {len(os.listdir(data_dir_woman))} images of woman.")
print(f"There are {len(os.listdir(data_dir_man))} images of man.")

```

There are 471 images of woman.

There are 472 images of man.

```

[7]: image_size = (img_size, img_size) # Resize images to this size

# Load train dataset (excluding validation & test set):
train_dataset = tf.keras.utils.image_dataset_from_directory(
    directory = DATA_DIR,
    image_size = image_size,
    batch_size = batch_size,
    label_mode='binary',
    validation_split = validation_split + test_split, # Total split for val +
↳test
    subset = "training",
    seed = seed_num
)

# Load validation dataset
val_dataset = tf.keras.utils.image_dataset_from_directory(
    directory = DATA_DIR,
    image_size = image_size,
    batch_size = batch_size,
    label_mode='binary',
    validation_split = validation_split + test_split,
    subset = "validation",

```

```

        seed = seed_num
    )

    # Further manually split validation dataset to extract test dataset
    val_batches = tf.data.experimental.cardinality(val_dataset)
    # Compute test dataset size (number of batches)
    test_size = round(val_batches.numpy() * (test_split / (validation_split +
        ↪test_split)))
    # Split validation dataset into validation and test subsets
    test_dataset = val_dataset.take(test_size)
    val_dataset = val_dataset.skip(test_size)

    print(f"Train batches: {tf.data.experimental.cardinality(train_dataset).
        ↪numpy()}")
    print(f"Validation batches: {tf.data.experimental.cardinality(val_dataset).
        ↪numpy()}")
    print(f"Test batches: {tf.data.experimental.cardinality(test_dataset).numpy()}")

    # Optimize for performance
    AUTOTUNE = tf.data.AUTOTUNE
    training_dataset = train_dataset.cache().shuffle(shuffle_buffer_size).
        ↪prefetch(buffer_size = AUTOTUNE)
    validation_dataset = val_dataset.cache().prefetch(buffer_size = AUTOTUNE)
    test_dataset = test_dataset.cache().prefetch(buffer_size = AUTOTUNE)

```

```

Found 943 files belonging to 2 classes.
Using 849 files for training.
Found 943 files belonging to 2 classes.
Using 94 files for validation.
Train batches: 27
Validation batches: 3
Test batches: 0

```

```

[8]: # Get the first batch of images and labels
for images, labels in training_dataset.take(1):
    example_batch_images = images
    example_batch_labels = labels

max_pixel = np.max(example_batch_images)
print(f"Maximum pixel value of images: {max_pixel}\n")
print(f"Shape of batch of images: {example_batch_images.shape}")
print(f"Shape of batch of labels: {example_batch_labels.shape}")

```

```

Maximum pixel value of images: 255.0

```

```

Shape of batch of images: (32, 130, 130, 3)
Shape of batch of labels: (32, 1)

```

```
[9]: '''
class EarlyStoppingCallback(tf.keras.callbacks.Callback):
    def on_epoch_end(self, epoch, logs=None):
        train_accuracy = logs.get('accuracy')
        val_accuracy = logs.get('val_accuracy')
        if train_accuracy >= desired_accuracy and val_accuracy >=
↪desired_accuracy:
            self.model.stop_training = True
            print(f"\nReached {desired_accuracy}% accuracy so cancelling
↪training!")
'''
```

```
[9]: '\n\nclass EarlyStoppingCallback(tf.keras.callbacks.Callback):\n    def
on_epoch_end(self, epoch, logs=None):\n        train_accuracy =
logs.get(\'accuracy\')\n        val_accuracy = logs.get(\'val_accuracy\')\n
if train_accuracy >= desired_accuracy and val_accuracy >= desired_accuracy:\n
self.model.stop_training = True\n        print(f"\nReached
{desired_accuracy}% accuracy so cancelling training!")\n'
```

```
[10]: '''
from tensorflow.keras.callbacks import EarlyStopping
early_stop = EarlyStopping(monitor='val_loss', patience=3)
'''
```

```
[10]: "\n\nfrom tensorflow.keras.callbacks import EarlyStopping\nearly_stop =
EarlyStopping(monitor='val_loss', patience=3)\n"
```

```
[11]: from tensorflow.keras.callbacks import LearningRateScheduler

# Reduce LR every 10 epochs (Learning rate decay factor)
def scheduler(epoch, lr):
    if epoch < 20:
        if epoch < 15:
            if epoch % 10 == 0 and epoch > 0:
                #return lr * 0.1
                return lr / 2
            return lr
        elif epoch % 10 == 0 and epoch > 0:
            #return lr * 0.1
            return lr / 5
        return lr
    elif epoch % 10 == 0 and epoch > 0:
        #return lr * 0.1
        return lr / 10
    return lr

lr_callback = LearningRateScheduler(scheduler)
```

```
[12]: # augmentation_model
def augment_model():
    """Creates a model (layers stacked on top of each other) for augmenting
    ↪ images of woman and man.

    Returns:
        tf.keras.Model: The model made up of the layers that will be used to
        ↪ augment the images of woman and man.
        """

    augmentation_model = tf.keras.Sequential([
        # Specify the input shape.
        tf.keras.Input(shape = (img_size, img_size, 3)),

        tf.keras.layers.RandomFlip("horizontal"),
        tf.keras.layers.RandomRotation(0.1, fill_mode = f_mode),
        #tf.keras.layers.RandomTranslation(0.1, 0.1, fill_mode = f_mode),
        #tf.keras.layers.RandomZoom(0.1, fill_mode=f_mode)
    ])

    return augmentation_model
```

```
[13]: def create_and_compile_model():
    """Creates, compiles and trains the model to predict woman and man images.

    Returns:
        tf.keras.Model: The model that will be trained to predict woman and man
        ↪ images.
        """

    augmentation_layers = augment_model()

    model = tf.keras.Sequential([
        # Note: the input shape is the desired size of the image: 150x150 with
        ↪ 3 bytes for color
        tf.keras.layers.InputLayer(shape = (img_size, img_size, 3)),
        augmentation_layers,
        tf.keras.layers.Rescaling(1./255),
        ##### CONV_LAYER_1: #####
        tf.keras.layers.Conv2D(n_filters_l1, (4, 4), activation = 'linear'),
        tf.keras.layers.MaxPooling2D(2, 2),
        ##### CONV_LAYER_2: #####
        tf.keras.layers.Conv2D(n_filters_l2, (3, 3), activation = 'relu'),
        tf.keras.layers.MaxPooling2D(2, 2),
        ##### CONV_LAYER_3: #####
        tf.keras.layers.Conv2D(64, (3, 3), activation = 'relu'),
        tf.keras.layers.MaxPooling2D(2, 2),
```

```

##### CONV_LAYER_4: #####
tf.keras.layers.Conv2D(64, (3, 3), activation = 'relu'),
tf.keras.layers.MaxPooling2D(2, 2),
##### CONV_LAYER_5: #####
tf.keras.layers.Conv2D(64, (3, 3), activation = 'relu'),
tf.keras.layers.MaxPooling2D(2, 2),
tf.keras.layers.Flatten(),
tf.keras.layers.Dropout(dropout_value),
##### BEFORE_LAST_LAYER: #####
tf.keras.layers.Dense(n_units_last_layer, activation = 'relu'),
# It will contain a value from 0-1 where 0 for the class 'female' and 1
↳ for the 'male'
tf.keras.layers.Dense(1, activation = 'sigmoid']]

model.compile(
    loss = loss,
    optimizer = optimizer,
    metrics = metrics
)

return model

```

```

[14]: # Create the compiled but untrained model
model = create_and_compile_model()
model.summary()

```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
sequential (Sequential)	(None, 130, 130, 3)	0
rescaling (Rescaling)	(None, 130, 130, 3)	0
conv2d (Conv2D)	(None, 127, 127, 16)	784
max_pooling2d (MaxPooling2D)	(None, 63, 63, 16)	0
conv2d_1 (Conv2D)	(None, 61, 61, 32)	4,640
max_pooling2d_1 (MaxPooling2D)	(None, 30, 30, 32)	0
conv2d_2 (Conv2D)	(None, 28, 28, 64)	18,496
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 64)	0

conv2d_3 (Conv2D)	(None, 12, 12, 64)	36,928
max_pooling2d_3 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_4 (Conv2D)	(None, 4, 4, 64)	36,928
max_pooling2d_4 (MaxPooling2D)	(None, 2, 2, 64)	0
flatten (Flatten)	(None, 256)	0
dropout (Dropout)	(None, 256)	0
dense (Dense)	(None, 2048)	526,336
dense_1 (Dense)	(None, 1)	2,049

Total params: 626,161 (2.39 MB)

Trainable params: 626,161 (2.39 MB)

Non-trainable params: 0 (0.00 B)

```
[15]: '''
training_history = model.fit(
    training_dataset,
    epochs = epochs,
    validation_data = validation_dataset,
    callbacks = [EarlyStoppingCallback()],
    verbose = 2
)
'''
```

```
[15]: '\ntraining_history = model.fit(\n    training_dataset,\n    epochs = epochs,\n    validation_data = validation_dataset,\n    callbacks =\n    [EarlyStoppingCallback()],\n    verbose = 2\n)\n'
```

```
[16]: '''
training_history = model.fit(
    training_dataset,
    epochs = epochs,
    validation_data = validation_dataset,
    callbacks=[early_stop],
    verbose = 2
)
'''
```

```
'''
```

```
[16]: '\ntraining_history = model.fit(\n    training_dataset,\n    epochs = epochs,\n    validation_data = validation_dataset,\n    callbacks=[early_stop],\n    verbose = 2\n)\n'
```

```
[17]: training_history = model.fit(\n    training_dataset,\n    epochs = epochs,\n    validation_data = validation_dataset,\n    callbacks = [lr_callback],\n    verbose = 2\n)
```

Epoch 1/100

27/27 - 3s - 116ms/step - accuracy: 0.4888 - loss: 0.6975 - val\_accuracy: 0.5638  
- val\_loss: 0.6900 - learning\_rate: 0.0010

Epoch 2/100

27/27 - 1s - 42ms/step - accuracy: 0.5406 - loss: 0.6868 - val\_accuracy: 0.4362  
- val\_loss: 0.7054 - learning\_rate: 0.0010

Epoch 3/100

27/27 - 1s - 43ms/step - accuracy: 0.6631 - loss: 0.6344 - val\_accuracy: 0.6702  
- val\_loss: 0.6337 - learning\_rate: 0.0010

Epoch 4/100

27/27 - 1s - 45ms/step - accuracy: 0.6949 - loss: 0.5823 - val\_accuracy: 0.6596  
- val\_loss: 0.6103 - learning\_rate: 0.0010

Epoch 5/100

27/27 - 1s - 44ms/step - accuracy: 0.7244 - loss: 0.5699 - val\_accuracy: 0.6702  
- val\_loss: 0.6025 - learning\_rate: 0.0010

Epoch 6/100

27/27 - 1s - 43ms/step - accuracy: 0.7550 - loss: 0.5296 - val\_accuracy: 0.7234  
- val\_loss: 0.5198 - learning\_rate: 0.0010

Epoch 7/100

27/27 - 1s - 43ms/step - accuracy: 0.7644 - loss: 0.4999 - val\_accuracy: 0.7660  
- val\_loss: 0.5145 - learning\_rate: 0.0010

Epoch 8/100

27/27 - 1s - 46ms/step - accuracy: 0.7585 - loss: 0.5229 - val\_accuracy: 0.7660  
- val\_loss: 0.5214 - learning\_rate: 0.0010

Epoch 9/100

27/27 - 1s - 46ms/step - accuracy: 0.7538 - loss: 0.4912 - val\_accuracy: 0.7872  
- val\_loss: 0.5634 - learning\_rate: 0.0010

Epoch 10/100

27/27 - 1s - 47ms/step - accuracy: 0.7927 - loss: 0.4641 - val\_accuracy: 0.8298  
- val\_loss: 0.4318 - learning\_rate: 0.0010

Epoch 11/100

27/27 - 1s - 45ms/step - accuracy: 0.8269 - loss: 0.4014 - val\_accuracy: 0.8511  
- val\_loss: 0.3682 - learning\_rate: 5.0000e-04

Epoch 12/100



27/27 - 1s - 46ms/step - accuracy: 0.8163 - loss: 0.4049 - val\_accuracy: 0.8404  
- val\_loss: 0.4428 - learning\_rate: 5.0000e-04  
Epoch 13/100  
27/27 - 1s - 43ms/step - accuracy: 0.8363 - loss: 0.3854 - val\_accuracy: 0.8511  
- val\_loss: 0.3520 - learning\_rate: 5.0000e-04  
Epoch 14/100  
27/27 - 1s - 43ms/step - accuracy: 0.8492 - loss: 0.3628 - val\_accuracy: 0.8511  
- val\_loss: 0.3256 - learning\_rate: 5.0000e-04  
Epoch 15/100  
27/27 - 1s - 43ms/step - accuracy: 0.8304 - loss: 0.3717 - val\_accuracy: 0.8936  
- val\_loss: 0.3001 - learning\_rate: 5.0000e-04  
Epoch 16/100  
27/27 - 1s - 43ms/step - accuracy: 0.8516 - loss: 0.3747 - val\_accuracy: 0.8511  
- val\_loss: 0.3405 - learning\_rate: 5.0000e-04  
Epoch 17/100  
27/27 - 1s - 43ms/step - accuracy: 0.8481 - loss: 0.3418 - val\_accuracy: 0.8617  
- val\_loss: 0.4021 - learning\_rate: 5.0000e-04  
Epoch 18/100  
27/27 - 1s - 45ms/step - accuracy: 0.8634 - loss: 0.3230 - val\_accuracy: 0.8617  
- val\_loss: 0.3437 - learning\_rate: 5.0000e-04  
Epoch 19/100  
27/27 - 1s - 43ms/step - accuracy: 0.8433 - loss: 0.3539 - val\_accuracy: 0.8830  
- val\_loss: 0.2870 - learning\_rate: 5.0000e-04  
Epoch 20/100  
27/27 - 1s - 42ms/step - accuracy: 0.8728 - loss: 0.3065 - val\_accuracy: 0.9043  
- val\_loss: 0.2947 - learning\_rate: 5.0000e-04  
Epoch 21/100  
27/27 - 1s - 46ms/step - accuracy: 0.8751 - loss: 0.2799 - val\_accuracy: 0.8723  
- val\_loss: 0.3039 - learning\_rate: 5.0000e-05  
Epoch 22/100  
27/27 - 1s - 46ms/step - accuracy: 0.8881 - loss: 0.2736 - val\_accuracy: 0.8617  
- val\_loss: 0.2974 - learning\_rate: 5.0000e-05  
Epoch 23/100  
27/27 - 1s - 49ms/step - accuracy: 0.8693 - loss: 0.2917 - val\_accuracy: 0.8830  
- val\_loss: 0.2915 - learning\_rate: 5.0000e-05  
Epoch 24/100  
27/27 - 1s - 46ms/step - accuracy: 0.8857 - loss: 0.2651 - val\_accuracy: 0.8936  
- val\_loss: 0.2894 - learning\_rate: 5.0000e-05  
Epoch 25/100  
27/27 - 1s - 43ms/step - accuracy: 0.8940 - loss: 0.2575 - val\_accuracy: 0.8617  
- val\_loss: 0.2972 - learning\_rate: 5.0000e-05  
Epoch 26/100  
27/27 - 1s - 43ms/step - accuracy: 0.8881 - loss: 0.2862 - val\_accuracy: 0.8936  
- val\_loss: 0.2774 - learning\_rate: 5.0000e-05  
Epoch 27/100  
27/27 - 1s - 44ms/step - accuracy: 0.8963 - loss: 0.2640 - val\_accuracy: 0.8830  
- val\_loss: 0.2864 - learning\_rate: 5.0000e-05  
Epoch 28/100

27/27 - 1s - 43ms/step - accuracy: 0.8905 - loss: 0.2519 - val\_accuracy: 0.8830  
- val\_loss: 0.2833 - learning\_rate: 5.0000e-05  
Epoch 29/100  
27/27 - 1s - 43ms/step - accuracy: 0.9034 - loss: 0.2560 - val\_accuracy: 0.8936  
- val\_loss: 0.2816 - learning\_rate: 5.0000e-05  
Epoch 30/100  
27/27 - 1s - 43ms/step - accuracy: 0.9034 - loss: 0.2554 - val\_accuracy: 0.8723  
- val\_loss: 0.2932 - learning\_rate: 5.0000e-05  
Epoch 31/100  
27/27 - 1s - 43ms/step - accuracy: 0.8987 - loss: 0.2553 - val\_accuracy: 0.8830  
- val\_loss: 0.2910 - learning\_rate: 5.0000e-06  
Epoch 32/100  
27/27 - 1s - 43ms/step - accuracy: 0.8857 - loss: 0.2610 - val\_accuracy: 0.8830  
- val\_loss: 0.2905 - learning\_rate: 5.0000e-06  
Epoch 33/100  
27/27 - 1s - 47ms/step - accuracy: 0.8940 - loss: 0.2570 - val\_accuracy: 0.8830  
- val\_loss: 0.2884 - learning\_rate: 5.0000e-06  
Epoch 34/100  
27/27 - 1s - 51ms/step - accuracy: 0.8893 - loss: 0.2586 - val\_accuracy: 0.8830  
- val\_loss: 0.2862 - learning\_rate: 5.0000e-06  
Epoch 35/100  
27/27 - 1s - 50ms/step - accuracy: 0.8893 - loss: 0.2669 - val\_accuracy: 0.8936  
- val\_loss: 0.2882 - learning\_rate: 5.0000e-06  
Epoch 36/100  
27/27 - 1s - 47ms/step - accuracy: 0.9011 - loss: 0.2416 - val\_accuracy: 0.8936  
- val\_loss: 0.2863 - learning\_rate: 5.0000e-06  
Epoch 37/100  
27/27 - 1s - 46ms/step - accuracy: 0.9011 - loss: 0.2408 - val\_accuracy: 0.8830  
- val\_loss: 0.2847 - learning\_rate: 5.0000e-06  
Epoch 38/100  
27/27 - 1s - 44ms/step - accuracy: 0.8975 - loss: 0.2581 - val\_accuracy: 0.8830  
- val\_loss: 0.2836 - learning\_rate: 5.0000e-06  
Epoch 39/100  
27/27 - 1s - 44ms/step - accuracy: 0.8822 - loss: 0.2641 - val\_accuracy: 0.8936  
- val\_loss: 0.2846 - learning\_rate: 5.0000e-06  
Epoch 40/100  
27/27 - 1s - 46ms/step - accuracy: 0.8869 - loss: 0.2693 - val\_accuracy: 0.8936  
- val\_loss: 0.2853 - learning\_rate: 5.0000e-06  
Epoch 41/100  
27/27 - 1s - 48ms/step - accuracy: 0.8822 - loss: 0.2707 - val\_accuracy: 0.8936  
- val\_loss: 0.2850 - learning\_rate: 5.0000e-07  
Epoch 42/100  
27/27 - 1s - 46ms/step - accuracy: 0.8893 - loss: 0.2517 - val\_accuracy: 0.8936  
- val\_loss: 0.2851 - learning\_rate: 5.0000e-07  
Epoch 43/100  
27/27 - 1s - 46ms/step - accuracy: 0.8822 - loss: 0.2510 - val\_accuracy: 0.8936  
- val\_loss: 0.2850 - learning\_rate: 5.0000e-07  
Epoch 44/100

27/27 - 1s - 49ms/step - accuracy: 0.8999 - loss: 0.2591 - val\_accuracy: 0.8936  
- val\_loss: 0.2848 - learning\_rate: 5.0000e-07  
Epoch 45/100  
27/27 - 1s - 48ms/step - accuracy: 0.8905 - loss: 0.2581 - val\_accuracy: 0.8936  
- val\_loss: 0.2847 - learning\_rate: 5.0000e-07  
Epoch 46/100  
27/27 - 1s - 46ms/step - accuracy: 0.8963 - loss: 0.2666 - val\_accuracy: 0.8936  
- val\_loss: 0.2846 - learning\_rate: 5.0000e-07  
Epoch 47/100  
27/27 - 1s - 48ms/step - accuracy: 0.8999 - loss: 0.2491 - val\_accuracy: 0.8936  
- val\_loss: 0.2846 - learning\_rate: 5.0000e-07  
Epoch 48/100  
27/27 - 1s - 51ms/step - accuracy: 0.8975 - loss: 0.2605 - val\_accuracy: 0.8936  
- val\_loss: 0.2844 - learning\_rate: 5.0000e-07  
Epoch 49/100  
27/27 - 1s - 50ms/step - accuracy: 0.8857 - loss: 0.2674 - val\_accuracy: 0.8936  
- val\_loss: 0.2842 - learning\_rate: 5.0000e-07  
Epoch 50/100  
27/27 - 1s - 48ms/step - accuracy: 0.9069 - loss: 0.2446 - val\_accuracy: 0.8936  
- val\_loss: 0.2842 - learning\_rate: 5.0000e-07  
Epoch 51/100  
27/27 - 1s - 45ms/step - accuracy: 0.8963 - loss: 0.2694 - val\_accuracy: 0.8936  
- val\_loss: 0.2841 - learning\_rate: 5.0000e-08  
Epoch 52/100  
27/27 - 1s - 46ms/step - accuracy: 0.8940 - loss: 0.2564 - val\_accuracy: 0.8936  
- val\_loss: 0.2841 - learning\_rate: 5.0000e-08  
Epoch 53/100  
27/27 - 1s - 45ms/step - accuracy: 0.8916 - loss: 0.2504 - val\_accuracy: 0.8936  
- val\_loss: 0.2841 - learning\_rate: 5.0000e-08  
Epoch 54/100  
27/27 - 1s - 46ms/step - accuracy: 0.8916 - loss: 0.2503 - val\_accuracy: 0.8936  
- val\_loss: 0.2841 - learning\_rate: 5.0000e-08  
Epoch 55/100  
27/27 - 1s - 46ms/step - accuracy: 0.9011 - loss: 0.2466 - val\_accuracy: 0.8936  
- val\_loss: 0.2841 - learning\_rate: 5.0000e-08  
Epoch 56/100  
27/27 - 1s - 45ms/step - accuracy: 0.8999 - loss: 0.2585 - val\_accuracy: 0.8936  
- val\_loss: 0.2841 - learning\_rate: 5.0000e-08  
Epoch 57/100  
27/27 - 1s - 45ms/step - accuracy: 0.9011 - loss: 0.2379 - val\_accuracy: 0.8936  
- val\_loss: 0.2841 - learning\_rate: 5.0000e-08  
Epoch 58/100  
27/27 - 1s - 46ms/step - accuracy: 0.8928 - loss: 0.2587 - val\_accuracy: 0.8936  
- val\_loss: 0.2841 - learning\_rate: 5.0000e-08  
Epoch 59/100  
27/27 - 1s - 46ms/step - accuracy: 0.8963 - loss: 0.2505 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-08  
Epoch 60/100

27/27 - 1s - 49ms/step - accuracy: 0.8905 - loss: 0.2524 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-08  
Epoch 61/100  
27/27 - 1s - 49ms/step - accuracy: 0.9022 - loss: 0.2532 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-09  
Epoch 62/100  
27/27 - 1s - 50ms/step - accuracy: 0.9022 - loss: 0.2580 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-09  
Epoch 63/100  
27/27 - 1s - 47ms/step - accuracy: 0.8905 - loss: 0.2535 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-09  
Epoch 64/100  
27/27 - 1s - 45ms/step - accuracy: 0.8928 - loss: 0.2421 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-09  
Epoch 65/100  
27/27 - 1s - 46ms/step - accuracy: 0.8940 - loss: 0.2467 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-09  
Epoch 66/100  
27/27 - 1s - 46ms/step - accuracy: 0.9022 - loss: 0.2371 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-09  
Epoch 67/100  
27/27 - 1s - 45ms/step - accuracy: 0.8975 - loss: 0.2512 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-09  
Epoch 68/100  
27/27 - 1s - 46ms/step - accuracy: 0.8963 - loss: 0.2506 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-09  
Epoch 69/100  
27/27 - 1s - 46ms/step - accuracy: 0.9022 - loss: 0.2312 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-09  
Epoch 70/100  
27/27 - 1s - 48ms/step - accuracy: 0.9128 - loss: 0.2462 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-09  
Epoch 71/100  
27/27 - 1s - 47ms/step - accuracy: 0.9081 - loss: 0.2426 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-10  
Epoch 72/100  
27/27 - 1s - 47ms/step - accuracy: 0.8940 - loss: 0.2706 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-10  
Epoch 73/100  
27/27 - 1s - 51ms/step - accuracy: 0.9058 - loss: 0.2472 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-10  
Epoch 74/100  
27/27 - 1s - 54ms/step - accuracy: 0.8940 - loss: 0.2527 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-10  
Epoch 75/100  
27/27 - 1s - 51ms/step - accuracy: 0.8893 - loss: 0.2606 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-10  
Epoch 76/100

27/27 - 1s - 46ms/step - accuracy: 0.9058 - loss: 0.2279 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-10  
Epoch 77/100  
27/27 - 1s - 46ms/step - accuracy: 0.8928 - loss: 0.2612 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-10  
Epoch 78/100  
27/27 - 1s - 48ms/step - accuracy: 0.8775 - loss: 0.2768 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-10  
Epoch 79/100  
27/27 - 1s - 46ms/step - accuracy: 0.8893 - loss: 0.2516 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-10  
Epoch 80/100  
27/27 - 1s - 46ms/step - accuracy: 0.9046 - loss: 0.2451 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-10  
Epoch 81/100  
27/27 - 1s - 46ms/step - accuracy: 0.8893 - loss: 0.2491 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-11  
Epoch 82/100  
27/27 - 1s - 46ms/step - accuracy: 0.8834 - loss: 0.2630 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-11  
Epoch 83/100  
27/27 - 1s - 46ms/step - accuracy: 0.8963 - loss: 0.2604 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-11  
Epoch 84/100  
27/27 - 1s - 45ms/step - accuracy: 0.8940 - loss: 0.2529 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-11  
Epoch 85/100  
27/27 - 1s - 50ms/step - accuracy: 0.8881 - loss: 0.2530 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-11  
Epoch 86/100  
27/27 - 1s - 49ms/step - accuracy: 0.8846 - loss: 0.2639 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-11  
Epoch 87/100  
27/27 - 1s - 48ms/step - accuracy: 0.8963 - loss: 0.2574 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-11  
Epoch 88/100  
27/27 - 1s - 48ms/step - accuracy: 0.8940 - loss: 0.2453 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-11  
Epoch 89/100  
27/27 - 1s - 45ms/step - accuracy: 0.8810 - loss: 0.2711 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-11  
Epoch 90/100  
27/27 - 1s - 45ms/step - accuracy: 0.8822 - loss: 0.2577 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-11  
Epoch 91/100  
27/27 - 1s - 45ms/step - accuracy: 0.9022 - loss: 0.2483 - val\_accuracy: 0.8936  
- val\_loss: 0.2840 - learning\_rate: 5.0000e-12  
Epoch 92/100

```

27/27 - 1s - 45ms/step - accuracy: 0.8881 - loss: 0.2500 - val_accuracy: 0.8936
- val_loss: 0.2840 - learning_rate: 5.0000e-12
Epoch 93/100
27/27 - 1s - 45ms/step - accuracy: 0.8940 - loss: 0.2529 - val_accuracy: 0.8936
- val_loss: 0.2840 - learning_rate: 5.0000e-12
Epoch 94/100
27/27 - 1s - 46ms/step - accuracy: 0.8940 - loss: 0.2391 - val_accuracy: 0.8936
- val_loss: 0.2840 - learning_rate: 5.0000e-12
Epoch 95/100
27/27 - 1s - 47ms/step - accuracy: 0.9022 - loss: 0.2502 - val_accuracy: 0.8936
- val_loss: 0.2840 - learning_rate: 5.0000e-12
Epoch 96/100
27/27 - 1s - 47ms/step - accuracy: 0.8928 - loss: 0.2543 - val_accuracy: 0.8936
- val_loss: 0.2840 - learning_rate: 5.0000e-12
Epoch 97/100
27/27 - 1s - 47ms/step - accuracy: 0.8928 - loss: 0.2577 - val_accuracy: 0.8936
- val_loss: 0.2840 - learning_rate: 5.0000e-12
Epoch 98/100
27/27 - 1s - 49ms/step - accuracy: 0.8810 - loss: 0.2772 - val_accuracy: 0.8936
- val_loss: 0.2840 - learning_rate: 5.0000e-12
Epoch 99/100
27/27 - 1s - 50ms/step - accuracy: 0.8881 - loss: 0.2635 - val_accuracy: 0.8936
- val_loss: 0.2840 - learning_rate: 5.0000e-12
Epoch 100/100
27/27 - 1s - 49ms/step - accuracy: 0.8893 - loss: 0.2637 - val_accuracy: 0.8936
- val_loss: 0.2840 - learning_rate: 5.0000e-12

```

```

[18]: #from tensorflow.keras.models import load_model
      #model.save('gender_recognition_project04_v10.h5')

```

```

[19]: model.metrics_names

```

```

[19]: ['loss', 'compile_metrics']

```

```

[20]: result_history = pd.DataFrame(model.history.history)
      result_history.head(15)

```

```


[20]:   accuracy    loss  val_accuracy  val_loss  learning_rate
0   0.488810  0.697535    0.563830  0.689982         0.0010
1   0.540636  0.686790    0.436170  0.705353         0.0010
2   0.663133  0.634382    0.670213  0.633700         0.0010
3   0.694935  0.582258    0.659574  0.610273         0.0010
4   0.724382  0.569885    0.670213  0.602472         0.0010
5   0.755006  0.529599    0.723404  0.519779         0.0010
6   0.764429  0.499851    0.765957  0.514465         0.0010
7   0.758539  0.522868    0.765957  0.521371         0.0010
8   0.753828  0.491202    0.787234  0.563364         0.0010
9   0.792697  0.464055    0.829787  0.431788         0.0010

```

10	0.826855	0.401430	0.851064	0.368193	0.0005
11	0.816254	0.404856	0.840426	0.442838	0.0005
12	0.836278	0.385424	0.851064	0.352016	0.0005
13	0.849234	0.362829	0.851064	0.325566	0.0005
14	0.830389	0.371654	0.893617	0.300061	0.0005


```
[21]: result_history[['loss', 'val_loss']].plot(figsize=(5, 3))
```

```
[21]: <Axes: >
```



```
[22]: result_history[['accuracy', 'val_accuracy']].plot(figsize=(5, 3))
```

```
[22]: <Axes: >
```



```
[23]: print(model.metrics_names)
      print(model.evaluate(validation_dataset))
```

```
['loss', 'compile_metrics']
3/3          0s 17ms/step -
accuracy: 0.8843 - loss: 0.2957
[0.2840394079685211, 0.8936170339584351]
```

```
[24]: from sklearn.metrics import classification_report, confusion_matrix

y_true = np.concatenate([y.numpy() for _, y in validation_dataset])
y_pred_prob = model.predict(validation_dataset)
# Convert probabilities to class labels (0:Female or 1:Male)
y_pred = (y_pred_prob > 0.5).astype(int).flatten()

print("Classification Report:\n", classification_report(y_true, y_pred,
    ↪target_names=['Female', 'Male']))
```

```
3/3          0s 47ms/step
Classification Report:
              precision    recall  f1-score   support

   Female       0.83        0.95        0.89         41
    Male       0.96        0.85        0.90         53

 accuracy                   0.89         94
 macro avg       0.89        0.90        0.89         94
weighted avg       0.90        0.89        0.89         94
```

```
[25]: import tensorflow as tf
      import numpy as np
      import matplotlib.pyplot as plt
      from tensorflow.keras.models import Model
      from tensorflow.keras.utils import load_img, img_to_array

img_size = img_size
model = tf.keras.models.load_model("gender_recognition_project04_v10.h5")

# Load your personal image if you are interested to predict:
your_image_path = "D:\\Hossein's desktop files in Microsoft Studio_
    ↪Laptop\\Personal Photos\\Hossein_10.jpg"

img = load_img(your_image_path, target_size=(img_size, img_size))
final_img = img_to_array(img)
# Adding a batch dimension:
```



```

final_img = np.expand_dims(final_img, axis=0)
prediction = model.predict(final_img)
result = "Female" if prediction > 0.5 else "Male"
if result=="Female":
    confidence = (model.predict(final_img)[0][0])*100
else:
    confidence = (1-model.predict(final_img)[0][0])*100
print(f"Prediction result: {result} (confidence= {confidence:.2f} %)")

# Visualize CNN Layers
successive_feature_maps = visualization_model.predict(final_img)
layer_names = [layer.name for layer in model.layers]

for layer_name, feature_map in zip(layer_names, successive_feature_maps):
    if len(feature_map.shape) == 4: # Only visualize conv/maxpool layers
        n_features = feature_map.shape[-1] # Number of filters
        size = feature_map.shape[1] # Feature map size
        display_grid = np.zeros((size, size * n_features))

        for i in range(n_features):
            x = feature_map[0, :, :, i]
            x -= x.mean()
            x /= (x.std() + 1e-8) # Normalize
            x *= 64
            x += 128
            x = np.clip(x, 0, 255).astype('uint8') # Convert to image format
            display_grid[:, i * size: (i + 1) * size] = x

        scale = 20. / n_features
        plt.figure(figsize=(scale * n_features, scale))
        plt.title(layer_name)
        plt.grid(False)
        plt.imshow(display_grid, aspect='auto', cmap='cividis')
        plt.show()

```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile\_metrics` will be empty until you train or evaluate the model.

```

1/1          0s 88ms/step
1/1          0s 46ms/step
Prediction result: Male (confidence= 94.19 %)

```

```

-----
NameError                                Traceback (most recent call last)
Cell In[25], line 26
    23 print(f"Prediction result: {result} (confidence= {confidence:.2f} %)")
    25 # Visualize CNN Layers

```

```
---> 26 successive_feature_maps = visualization_model.predict(final_img)
      27 layer_names = [layer.name for layer in model.layers]
      29 for layer_name, feature_map in zip(layer_names, successive_feature_maps :
```

```
NameError: name 'visualization_model' is not defined
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```