

Personal Project_04_v10_test1_4conv-layer_run78_very advanced control 4_autorun

May 7, 2025

```
[1]: from tensorflow.keras.callbacks import LearningRateScheduler
from sklearn.metrics import classification_report, confusion_matrix
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import matplotlib.image as mpimg
import tensorflow as tf
import os

class EarlyStoppingCallback(tf.keras.callbacks.Callback):
    def on_epoch_end(self, epoch, logs=None):
        train_accuracy = logs.get('accuracy')
        val_accuracy = logs.get('val_accuracy')
        if train_accuracy >= desired_train_accuracy and val_accuracy >=
desired_val_accuracy:
            self.model.stop_training = True
            print("Reached desired accuracy so cancelling training!")

# target accuracy values:
desired_train_accuracy = 0.89
desired_val_accuracy = 0.89

# maximum trial number:
trial_num = 50

# maximum possible epoch:
epochs = 20

TRAIN_ACC=0.1
VAL_ACC=0.1
try_num = 1
condition = True
```

```

while (try_num<trial_num and condition==True):
    # DOE factors:
    learning_rate = 0.0005
    dropout_value = 0.2
    # n-conv_layers = 4
    n_units_last_layer = 4096
    n_filters_l1 = 32
    n_filters_l2 = 64

    # other factors:
    img_size = 130
    batch_size = 32
    validation_split = 0.1 # 10% for validation
    test_split = 0.00 # 0% for testing
    shuffle_buffer_size = 1000
    seed_num = 101
    desired_accuracy = 0.99 # it should be active if EarlyStoppingCallback is
    ↪activated
    loss = 'binary_crossentropy'
    #optimizer = tf.keras.optimizers.RMSprop(learning_rate=learning_rate)
    optimizer = tf.keras.optimizers.Adam(learning_rate=learning_rate)
    metrics = ['accuracy']
    f_mode = 'nearest' # fill_mode in image augmentation

    #DATA_DIR = "D:\\CS online courses\\Free DataSets\\Free Images\\Easier
    ↪portrait images_GPU_03"
    DATA_DIR = "/Users/hosseini/Downloads/Easier portrait images_GPU_03"

    # Subdirectories for each class
    data_dir_woman = os.path.join(DATA_DIR, 'woman')
    data_dir_man = os.path.join(DATA_DIR, 'man')

    image_size = (img_size, img_size) # Resize images to this size

    # Load train dataset (excluding validation & test set):
    train_dataset = tf.keras.utils.image_dataset_from_directory(
        directory = DATA_DIR,
        image_size = image_size,
        batch_size = batch_size,
        label_mode='binary',
        validation_split = validation_split + test_split, # Total split for
    ↪val + test
        subset = "training",

```

```

        seed = seed_num
    )

    # Load validation dataset
    val_dataset = tf.keras.utils.image_dataset_from_directory(
        directory = DATA_DIR,
        image_size = image_size,
        batch_size = batch_size,
        label_mode='binary',
        validation_split = validation_split + test_split,
        subset = "validation",
        seed = seed_num
    )

    # Further manually split validation dataset to extract test dataset
    val_batches = tf.data.experimental.cardinality(val_dataset)
    # Compute test dataset size (number of batches)
    test_size = round(val_batches.numpy() * (test_split / (validation_split +
↪test_split)))
    # Split validation dataset into validation and test subsets
    test_dataset = val_dataset.take(test_size)
    val_dataset = val_dataset.skip(test_size)

    # Optimize for performance
    AUTOTUNE = tf.data.AUTOTUNE
    training_dataset = train_dataset.cache().shuffle(shuffle_buffer_size).
↪prefetch(buffer_size = AUTOTUNE)
    validation_dataset = val_dataset.cache().prefetch(buffer_size = AUTOTUNE)
    test_dataset = test_dataset.cache().prefetch(buffer_size = AUTOTUNE)

    # Get the first batch of images and labels
    for images, labels in training_dataset.take(1):
        example_batch_images = images
        example_batch_labels = labels

    max_pixel = np.max(example_batch_images)

    def scheduler(epoch, lr):
        if epoch < 10:
            if epoch % 5 == 0 and epoch > 0:
                return lr / 1
            return lr
        elif epoch < 15:
            if epoch % 5 == 0 and epoch > 0:
                return lr / 2
            return lr
        elif epoch < 30:

```

```

        if epoch % 5 == 0 and epoch > 0:
            return lr / 1
        return lr
    return lr
lr_callback = LearningRateScheduler(scheduler)

# augmentation_model
def augment_model():
    augmentation_model = tf.keras.Sequential([
        # Specify the input shape.
        tf.keras.Input(shape = (img_size, img_size, 3)),
        tf.keras.layers.RandomFlip("horizontal"),
        tf.keras.layers.RandomRotation(0.1, fill_mode = f_mode),
        #tf.keras.layers.RandomTranslation(0.1, 0.1, fill_mode = f_mode),
        #tf.keras.layers.RandomZoom(0.1, fill_mode=f_mode)
    ])
    return augmentation_model

def create_and_compile_model():
    augmentation_layers = augment_model()
    model = tf.keras.Sequential([
        # Note: the input shape is the desired size of the image: 150x150
        ↪with 3 bytes for color
        tf.keras.layers.InputLayer(shape = (img_size, img_size, 3)),
        augmentation_layers,
        tf.keras.layers.Rescaling(1./255),
        ##### CONV_LAYER_1: #####
        tf.keras.layers.Conv2D(n_filters_l1, (4, 4), activation = 'linear'),
        tf.keras.layers.MaxPooling2D(2, 2),
        ##### CONV_LAYER_2: #####
        tf.keras.layers.Conv2D(n_filters_l2, (3, 3), activation = 'relu'),
        tf.keras.layers.MaxPooling2D(2, 2),
        ##### CONV_LAYER_3: #####
        tf.keras.layers.Conv2D(64, (3, 3), activation = 'relu'),
        tf.keras.layers.MaxPooling2D(2, 2),
        ##### CONV_LAYER_4: #####
        tf.keras.layers.Conv2D(64, (3, 3), activation = 'relu'),
        tf.keras.layers.MaxPooling2D(2, 2),
        tf.keras.layers.Flatten(),
        tf.keras.layers.Dropout(dropout_value),
        ##### BEFORE_LAST_LAYER: #####
        tf.keras.layers.Dense(n_units_last_layer, activation = 'relu'),
        # It will contain a value from 0-1 where 0 for the class 'female'
        ↪and 1 for the 'male'
        tf.keras.layers.Dense(1, activation = 'sigmoid'))
    model.compile(
        loss = loss,

```

```

        optimizer = optimizer,
        metrics = metrics
    )
    return model

# Create the compiled but untrained model
def reset_weights(model):
    for layer in model.layers:
        if hasattr(layer, 'kernel_initializer'):
            layer.kernel.assign(layer.kernel_initializer(layer.kernel.
↳shape))
        if hasattr(layer, 'bias_initializer'):
            layer.bias.assign(layer.bias_initializer(layer.bias.shape))

model = create_and_compile_model()
reset_weights(model) # Reset all layer weights
training_history = model.fit(training_dataset,
                             epochs=epochs,
                             validation_data=validation_dataset,
                             callbacks=[lr_callback,
↳EarlyStoppingCallback()],
                             verbose=1)
result_history = pd.DataFrame(model.history.history)
TRAIN_ACC = result_history['accuracy'].iloc[-1]
print(f"Current training accuracy: {TRAIN_ACC}")
VAL_ACC = result_history['val_accuracy'].iloc[-1]
print(f"Current validation accuracy: {VAL_ACC}")
# Restart script
print("Resetting all weights...")
print(f'Current number of trials: {try_num}')
try_num += 1
result_history[['loss', 'val_loss']].plot(figsize=(5, 3))
result_history[['accuracy', 'val_accuracy']].plot(figsize=(5, 3))
plt.show()
print(model.metrics_names)
print(model.evaluate(validation_dataset))
y_true = np.concatenate([y.numpy() for _, y in validation_dataset])
y_pred_prob = model.predict(validation_dataset)
# Convert probabilities to class labels (0:Female or 1:Male)
y_pred = (y_pred_prob > 0.5).astype(int).flatten()
print("Classification Report:\n", classification_report(y_true, y_pred,
↳target_names=['Female', 'Male']))

if (TRAIN_ACC>=desired_train_accuracy and VAL_ACC>=desired_val_accuracy):
    condition = False
    model.save('trained_model_run78_very_advanced_control.h5')

```

```
result_history.head(15)
```

Found 943 files belonging to 2 classes.

Using 849 files for training.

Found 943 files belonging to 2 classes.

Using 94 files for validation.

Epoch 1/20

2025-05-07 12:22:49.495620: I tensorflow/core/framework/local_rendezvous.cc:405]

Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence

27/27 3s 98ms/step -

accuracy: 0.5822 - loss: 0.6885 - val_accuracy: 0.5000 - val_loss: 0.8843 -

learning_rate: 5.0000e-04

Epoch 2/20

27/27 3s 95ms/step -

accuracy: 0.6371 - loss: 0.6461 - val_accuracy: 0.7234 - val_loss: 0.5644 -

learning_rate: 5.0000e-04

Epoch 3/20

27/27 3s 98ms/step -

accuracy: 0.7372 - loss: 0.5271 - val_accuracy: 0.7021 - val_loss: 0.6582 -

learning_rate: 5.0000e-04

Epoch 4/20

27/27 3s 95ms/step -

accuracy: 0.7479 - loss: 0.5298 - val_accuracy: 0.8191 - val_loss: 0.4458 -

learning_rate: 5.0000e-04

Epoch 5/20

27/27 3s 97ms/step -

accuracy: 0.7642 - loss: 0.4580 - val_accuracy: 0.8085 - val_loss: 0.4392 -

learning_rate: 5.0000e-04

Epoch 6/20

27/27 3s 96ms/step -

accuracy: 0.7932 - loss: 0.4415 - val_accuracy: 0.8511 - val_loss: 0.3819 -

learning_rate: 5.0000e-04

Epoch 7/20

27/27 3s 97ms/step -

accuracy: 0.8219 - loss: 0.4030 - val_accuracy: 0.8085 - val_loss: 0.4056 -

learning_rate: 5.0000e-04

Epoch 8/20

27/27 3s 96ms/step -

accuracy: 0.8265 - loss: 0.4040 - val_accuracy: 0.8085 - val_loss: 0.4204 -

learning_rate: 5.0000e-04

Epoch 9/20

27/27 3s 96ms/step -

accuracy: 0.8598 - loss: 0.3465 - val_accuracy: 0.8617 - val_loss: 0.4194 -



learning_rate: 5.0000e-04

Epoch 10/20

```

27/27          3s 96ms/step -
accuracy: 0.8604 - loss: 0.3343 - val_accuracy: 0.8723 - val_loss: 0.2692 -
learning_rate: 5.0000e-04
Epoch 11/20
27/27          3s 97ms/step -
accuracy: 0.8968 - loss: 0.2748 - val_accuracy: 0.8723 - val_loss: 0.4017 -
learning_rate: 2.5000e-04
Epoch 12/20
27/27          3s 96ms/step -
accuracy: 0.8852 - loss: 0.2814 - val_accuracy: 0.8404 - val_loss: 0.4579 -
learning_rate: 2.5000e-04
Epoch 13/20
27/27          3s 96ms/step -
accuracy: 0.8547 - loss: 0.3136 - val_accuracy: 0.8404 - val_loss: 0.4032 -
learning_rate: 2.5000e-04
Epoch 14/20
27/27          3s 96ms/step -
accuracy: 0.9058 - loss: 0.2361 - val_accuracy: 0.8723 - val_loss: 0.3769 -
learning_rate: 2.5000e-04
Epoch 15/20
27/27          3s 97ms/step -
accuracy: 0.8999 - loss: 0.2315 - val_accuracy: 0.8511 - val_loss: 0.3192 -
learning_rate: 2.5000e-04
Epoch 16/20
27/27          3s 97ms/step -
accuracy: 0.9107 - loss: 0.2446 - val_accuracy: 0.8511 - val_loss: 0.3997 -
learning_rate: 2.5000e-04
Epoch 17/20
27/27          3s 98ms/step -
accuracy: 0.9136 - loss: 0.2122 - val_accuracy: 0.8830 - val_loss: 0.3810 -
learning_rate: 2.5000e-04
Epoch 18/20
27/27          3s 97ms/step -
accuracy: 0.9086 - loss: 0.2189 - val_accuracy: 0.8830 - val_loss: 0.4060 -
learning_rate: 2.5000e-04
Epoch 19/20
27/27          3s 97ms/step -
accuracy: 0.9036 - loss: 0.2123 - val_accuracy: 0.8830 - val_loss: 0.4281 -
learning_rate: 2.5000e-04
Epoch 20/20
27/27          3s 97ms/step -
accuracy: 0.9264 - loss: 0.1941 - val_accuracy: 0.8511 - val_loss: 0.3769 -
learning_rate: 2.5000e-04
Current training accuracy: 0.9269729256629944
Current validation accuracy: 0.8510638475418091
Reseting all weights...
Current number of trials: 1

```



```
['loss', 'compile_metrics']
```

```
3/3          0s 25ms/step -
```

```
accuracy: 0.8318 - loss: 0.3831
```

```
[0.37694793939590454, 0.8510638475418091]
```

```
3/3          0s 36ms/step
```

```
Classification Report:
```

	precision	recall	f1-score	support
Female	0.81	0.85	0.83	41
Male	0.88	0.85	0.87	53
accuracy			0.85	94

macro avg	0.85	0.85	0.85	94
weighted avg	0.85	0.85	0.85	94

Found 943 files belonging to 2 classes.

Using 849 files for training.

Found 943 files belonging to 2 classes.

Using 94 files for validation.

2025-05-07 12:23:42.466022: I tensorflow/core/framework/local_rendezvous.cc:405]
Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence

Epoch 1/20

27/27 3s 100ms/step -

accuracy: 0.5201 - loss: 0.7306 - val_accuracy: 0.6596 - val_loss: 0.6441 -

learning_rate: 5.0000e-04

Epoch 2/20

27/27 3s 96ms/step -

accuracy: 0.6903 - loss: 0.5974 - val_accuracy: 0.7234 - val_loss: 0.5745 -

learning_rate: 5.0000e-04

Epoch 3/20

27/27 3s 98ms/step -

accuracy: 0.7712 - loss: 0.5243 - val_accuracy: 0.7234 - val_loss: 0.5244 -

learning_rate: 5.0000e-04

Epoch 4/20

27/27 3s 96ms/step -

accuracy: 0.7126 - loss: 0.5493 - val_accuracy: 0.7340 - val_loss: 0.5952 -

learning_rate: 5.0000e-04

Epoch 5/20

27/27 3s 96ms/step -

accuracy: 0.7730 - loss: 0.4587 - val_accuracy: 0.8191 - val_loss: 0.4628 -

learning_rate: 5.0000e-04

Epoch 6/20

27/27 3s 96ms/step -

accuracy: 0.7853 - loss: 0.4603 - val_accuracy: 0.7660 - val_loss: 0.5203 -

learning_rate: 5.0000e-04

Epoch 7/20

27/27 3s 96ms/step -

accuracy: 0.7849 - loss: 0.4686 - val_accuracy: 0.7553 - val_loss: 0.4762 -

learning_rate: 5.0000e-04

Epoch 8/20

27/27 3s 96ms/step -

accuracy: 0.7938 - loss: 0.4376 - val_accuracy: 0.7872 - val_loss: 0.4229 -

learning_rate: 5.0000e-04

Epoch 9/20

27/27 3s 100ms/step -

accuracy: 0.8030 - loss: 0.4236 - val_accuracy: 0.8191 - val_loss: 0.4136 -

learning_rate: 5.0000e-04



Epoch 10/20

27/27 3s 96ms/step -

```

accuracy: 0.8022 - loss: 0.4394 - val_accuracy: 0.8085 - val_loss: 0.4412 -
learning_rate: 5.0000e-04
Epoch 11/20
27/27          3s 97ms/step -
accuracy: 0.8375 - loss: 0.3721 - val_accuracy: 0.8617 - val_loss: 0.4235 -
learning_rate: 2.5000e-04
Epoch 12/20
27/27          3s 96ms/step -
accuracy: 0.8414 - loss: 0.3441 - val_accuracy: 0.8723 - val_loss: 0.3752 -
learning_rate: 2.5000e-04
Epoch 13/20
27/27          3s 96ms/step -
accuracy: 0.8530 - loss: 0.3435 - val_accuracy: 0.8404 - val_loss: 0.4519 -
learning_rate: 2.5000e-04
Epoch 14/20
27/27          3s 96ms/step -
accuracy: 0.8327 - loss: 0.3620 - val_accuracy: 0.8298 - val_loss: 0.4735 -
learning_rate: 2.5000e-04
Epoch 15/20
27/27          3s 98ms/step -
accuracy: 0.8462 - loss: 0.3397 - val_accuracy: 0.8511 - val_loss: 0.3953 -
learning_rate: 2.5000e-04
Epoch 16/20
27/27          3s 97ms/step -
accuracy: 0.8662 - loss: 0.2780 - val_accuracy: 0.8191 - val_loss: 0.3644 -
learning_rate: 2.5000e-04
Epoch 17/20
27/27          3s 97ms/step -
accuracy: 0.8829 - loss: 0.3070 - val_accuracy: 0.8617 - val_loss: 0.3110 -
learning_rate: 2.5000e-04
Epoch 18/20
27/27          3s 97ms/step -
accuracy: 0.8687 - loss: 0.2847 - val_accuracy: 0.8404 - val_loss: 0.3064 -
learning_rate: 2.5000e-04
Epoch 19/20
27/27          3s 97ms/step -
accuracy: 0.8816 - loss: 0.2991 - val_accuracy: 0.8830 - val_loss: 0.2557 -
learning_rate: 2.5000e-04
Epoch 20/20
27/27          3s 97ms/step -
accuracy: 0.8928 - loss: 0.2598 - val_accuracy: 0.8404 - val_loss: 0.3236 -
learning_rate: 2.5000e-04
Current training accuracy: 0.8951708078384399
Current validation accuracy: 0.8404255509376526
Reseting all weights...
Current number of trials: 2

```



```
['loss', 'compile_metrics']
3/3          0s 24ms/step -
accuracy: 0.8343 - loss: 0.3530
[0.32360365986824036, 0.8404255509376526]
```

```
3/3          0s 37ms/step
```

Classification Report:

	precision	recall	f1-score	support
Female	0.78	0.88	0.83	41
Male	0.90	0.81	0.85	53
accuracy			0.84	94

macro avg	0.84	0.84	0.84	94
weighted avg	0.85	0.84	0.84	94

Found 943 files belonging to 2 classes.

Using 849 files for training.

Found 943 files belonging to 2 classes.

Using 94 files for validation.

2025-05-07 12:24:35.960210: I tensorflow/core/framework/local_rendezvous.cc:405]
Local rendezvous is aborting with status: OUT_OF_RANGE: End of sequence

Epoch 1/20

27/27 3s 102ms/step -

accuracy: 0.5591 - loss: 0.7182 - val_accuracy: 0.7340 - val_loss: 0.6327 -

learning_rate: 5.0000e-04

Epoch 2/20

27/27 3s 97ms/step -

accuracy: 0.7219 - loss: 0.5909 - val_accuracy: 0.6809 - val_loss: 0.5680 -

learning_rate: 5.0000e-04

Epoch 3/20

27/27 3s 97ms/step -

accuracy: 0.6785 - loss: 0.5778 - val_accuracy: 0.7766 - val_loss: 0.4832 -

learning_rate: 5.0000e-04

Epoch 4/20

27/27 3s 97ms/step -

accuracy: 0.7385 - loss: 0.5046 - val_accuracy: 0.7872 - val_loss: 0.5345 -

learning_rate: 5.0000e-04

Epoch 5/20

27/27 3s 97ms/step -

accuracy: 0.7804 - loss: 0.4804 - val_accuracy: 0.8298 - val_loss: 0.4290 -

learning_rate: 5.0000e-04

Epoch 6/20

27/27 3s 98ms/step -

accuracy: 0.7826 - loss: 0.4872 - val_accuracy: 0.7979 - val_loss: 0.5635 -

learning_rate: 5.0000e-04

Epoch 7/20

27/27 3s 98ms/step -

accuracy: 0.7955 - loss: 0.4416 - val_accuracy: 0.8085 - val_loss: 0.3463 -

learning_rate: 5.0000e-04

Epoch 8/20

27/27 3s 97ms/step -

accuracy: 0.8053 - loss: 0.4130 - val_accuracy: 0.8298 - val_loss: 0.3492 -

learning_rate: 5.0000e-04

Epoch 9/20

27/27 3s 97ms/step -

accuracy: 0.8544 - loss: 0.3677 - val_accuracy: 0.8404 - val_loss: 0.3506 -

learning_rate: 5.0000e-04



Epoch 10/20

27/27 3s 97ms/step -

```

accuracy: 0.8611 - loss: 0.3603 - val_accuracy: 0.8617 - val_loss: 0.2943 -
learning_rate: 5.0000e-04
Epoch 11/20
27/27          3s 97ms/step -
accuracy: 0.8438 - loss: 0.3434 - val_accuracy: 0.8830 - val_loss: 0.3348 -
learning_rate: 2.5000e-04
Epoch 12/20
27/27          3s 98ms/step -
accuracy: 0.8594 - loss: 0.3544 - val_accuracy: 0.8723 - val_loss: 0.2727 -
learning_rate: 2.5000e-04
Epoch 13/20
27/27          3s 99ms/step -
accuracy: 0.8869 - loss: 0.2528 - val_accuracy: 0.8723 - val_loss: 0.3023 -
learning_rate: 2.5000e-04
Epoch 14/20
27/27          3s 98ms/step -
accuracy: 0.8975 - loss: 0.2524 - val_accuracy: 0.8511 - val_loss: 0.3737 -
learning_rate: 2.5000e-04
Epoch 15/20
27/27          3s 97ms/step -
accuracy: 0.8742 - loss: 0.2777 - val_accuracy: 0.8617 - val_loss: 0.3205 -
learning_rate: 2.5000e-04
Epoch 16/20
27/27          3s 100ms/step -
accuracy: 0.9109 - loss: 0.2373 - val_accuracy: 0.8830 - val_loss: 0.2283 -
learning_rate: 2.5000e-04
Epoch 17/20
27/27          0s 95ms/step -
accuracy: 0.9242 - loss: 0.2151Reached desired accuracy so cancelling training!
27/27          3s 98ms/step -
accuracy: 0.9233 - loss: 0.2161 - val_accuracy: 0.9149 - val_loss: 0.2582 -
learning_rate: 2.5000e-04
Current training accuracy: 0.8975265026092529
Current validation accuracy: 0.914893627166748
Reseting all weights...
Current number of trials: 3

```



```
['loss', 'compile_metrics']
```

```
3/3          0s 25ms/step -
```

```
accuracy: 0.8989 - loss: 0.2818
```

```
[0.25818532705307007, 0.914893627166748]
```

```
WARNING:tensorflow:5 out of the last 7 calls to <function
```

```
TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at
```

```
0x17f228fe0> triggered tf.function retracing. Tracing is expensive and the
```

```
excessive number of tracings could be due to (1) creating @tf.function
```

```
repeatedly in a loop, (2) passing tensors with different shapes, (3) passing
```

```
Python objects instead of tensors. For (1), please define your @tf.function
```

```
outside of the loop. For (2), @tf.function has reduce_retracing=True option that
```

```
can avoid unnecessary retracing. For (3), please refer to
```

https://www.tensorflow.org/guide/function#controlling_retracing and
https://www.tensorflow.org/api_docs/python/tf/function for more details.

```
1/3          0s
47ms/stepWARNING:tensorflow:6 out of the last 9 calls to <function
TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at
0x17f228fe0> triggered tf.function retracing. Tracing is expensive and the
excessive number of tracings could be due to (1) creating @tf.function
repeatedly in a loop, (2) passing tensors with different shapes, (3) passing
Python objects instead of tensors. For (1), please define your @tf.function
outside of the loop. For (2), @tf.function has reduce_retracing=True option that
can avoid unnecessary retracing. For (3), please refer to
https://www.tensorflow.org/guide/function#controlling\_retracing and
https://www.tensorflow.org/api\_docs/python/tf/function for more details.
3/3          0s 37ms/step
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or
`keras.saving.save_model(model)`. This file format is considered legacy. We
recommend using instead the native Keras format, e.g.
`model.save('my_model.keras')` or `keras.saving.save_model(model,
'my_model.keras')`.

Classification Report:

	precision	recall	f1-score	support
Female	0.85	0.98	0.91	41
Male	0.98	0.87	0.92	53
accuracy			0.91	94
macro avg	0.91	0.92	0.91	94
weighted avg	0.92	0.91	0.92	94

```
[1]: accuracy      loss  val_accuracy  val_loss  learning_rate
0    0.591284  0.700387    0.734043  0.632742    0.00050
1    0.722026  0.578647    0.680851  0.568016    0.00050
2    0.707892  0.555076    0.776596  0.483235    0.00050
3    0.740872  0.512653    0.787234  0.534464    0.00050
4    0.782097  0.474838    0.829787  0.429026    0.00050
5    0.786808  0.470764    0.797872  0.563464    0.00050
6    0.799764  0.429010    0.808511  0.346285    0.00050
7    0.815077  0.410316    0.829787  0.349219    0.00050
8    0.839812  0.380174    0.840426  0.350584    0.00050
9    0.851590  0.364124    0.861702  0.294300    0.00050
10   0.872792  0.296062    0.882979  0.334795    0.00025
11   0.855124  0.327675    0.872340  0.272722    0.00025
12   0.886926  0.270424    0.872340  0.302330    0.00025
13   0.888104  0.276107    0.851064  0.373744    0.00025
14   0.891637  0.257347    0.861702  0.320483    0.00025
```

```

[2]: from tensorflow.keras.models import Model
from tensorflow.keras.utils import load_img, img_to_array

img_size = img_size
model = tf.keras.models.load_model("trained_model_run78_very_advanced_control.
↳h5")

# Load your personal image if you are interested to predict:
#your_image_path = "D:\Hossein's desktop files in Microsoft Studio\
↳Laptop\Personal Photos\Hossein_10.jpg"
your_image_path = "/Users/hossein/Downloads/Hossein.png"

img = load_img(your_image_path, target_size=(img_size, img_size))
final_img = img_to_array(img)
# Adding a batch dimension:
final_img = np.expand_dims(final_img, axis=0)
prediction = model.predict(final_img)
result = "Female" if prediction > 0.5 else "Male"
if result=="Female":
    confidence = (model.predict(final_img)[0][0])*100
else:
    confidence = (1-model.predict(final_img)[0][0])*100
print(f"Prediction result: {result} (confidence= {confidence:.2f} %)")

# Visualize CNN Layers
dummy_input = np.random.rand(1, img_size, img_size, 3) # Create random input
↳with correct shape
model.predict(dummy_input) # Call the model to establish input tensors
visualization_model = Model(inputs=model.inputs, outputs=[layer.output for
↳layer in model.layers])
successive_feature_maps = visualization_model.predict(final_img)
layer_names = [layer.name for layer in model.layers]

for layer_name, feature_map in zip(layer_names, successive_feature_maps):
    if len(feature_map.shape) == 4: # Only visualize conv/maxpool layers
        n_features = feature_map.shape[-1] # Number of filters
        size = feature_map.shape[1] # Feature map size
        display_grid = np.zeros((size, size * n_features))

        for i in range(n_features):
            x = feature_map[0, :, :, i]
            x -= x.mean()
            x /= (x.std() + 1e-8) # Normalize
            x *= 64
            x += 128
            x = np.clip(x, 0, 255).astype('uint8') # Convert to image format
            display_grid[:, i * size: (i + 1) * size] = x

```



```

scale = 20. / n_features
plt.figure(figsize=(scale * n_features, scale))
plt.title(layer_name)
plt.grid(False)
plt.imshow(display_grid, aspect='auto', cmap='cividis')
plt.show()

```

WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the model.

1/1 0s 38ms/step

1/1 0s 16ms/step

Prediction result: Male (confidence= 91.61 %)

1/1 0s 15ms/step

/opt/anaconda3/envs/mytfenv/lib/python3.12/site-packages/keras/src/models/functional.py:237: UserWarning: The structure of `inputs` doesn't match the expected structure.






Expected: ['input_layer_5']


Received: inputs=Tensor(shape=(1, 130, 130, 3))

warnings.warn(msg)

1/1 0s 150ms/step







[]:

[]:

[]: